

Universitetet i Bergen  
Det matematisk-naturvitenskaplige fakultet

Exam in INF236 - Parallel programming,  
Monday 28.09.18, 9AM - 12AM  
English text

**Problem 1, speedup**

- A. A sequential algorithm spends  $1/4$  of its running time in an initial preprocessing phase, and  $1/4$  is spent in a postprocessing phase. Neither the preprocessing nor the postprocessing phase can be parallelized. What is the maximum speedup attainable when moving to a parallel version of the algorithm?

Let A, B, C and D be four programs that each take one minute to execute sequentially and that have to be run in sequence. Thus A must finish before B can start and so on. Assume that porting your programs to a GPU will give A a speedup of 100, B a speedup of 10, C a speedup of 1 and D a speedup of 0.1.

- B. What is the combined running time of the programs after porting all of them onto the GPU?
- C. Which programs would you recommend to move to the GPU and what is the combined running time of your solution.

**Problem 2, GPU questions**

Explain what must be done to synchronize the following groups of threads:

- A.
- I. All the threads in a warp
  - II. All the threads in a block
  - III. All the threads in the GPU

B.

Let A and T be two matrices of dimension  $n \times n$ . Then T is the *transpose* of A if  $T_{i,j} = A_{j,i}$  for every pair of values i and j. In the following you are to write a CUDA kernel for computing the transpose T of a matrix A. In the following kernel the variables A and T points to the first elements of the two matrices each containing  $n^2$  elements. Comment on if there are any restrictions on the block and grid dimension that your kernel can be called with.

```
__global__ void transpose(int *A,int *T,int n) {  
    // Your code goes here  
}
```

### Problem 3, BSP, sorting

Let A be an array of size n containing distinct floating point values (float). Thus A[i] and A[j] have different values for all i and j where i is not equal to j. We now want to sort the elements of A and store them in an array B, also of size n. We do this by counting the number of elements that are smaller than each A[i] and then placing A[i] in the correct position of B. The sequential code for this algorithm is as follows:

```
for (int i=0; i<n; i++) {
    int pos=0;
    for (int j=0; j<n; j++)
        if (A[j] < A[i]) pos++;
    B[pos] = a[i];
}
```

We now want to write a parallel algorithm using BSP for executing this algorithm on a distributed memory parallel computer. Initially every process has a part of the input of length np stored in a local array A. When the program is done each process should have np elements of the output stored in B. The elements of B on each processor should be sorted and all elements of B on process i should be smaller than all elements on process j if  $i < j$ .

The algorithm works by cyclicly shifting the elements of A between the processes (similarly to what was done in the first part of the obligatory assignment on matrix multiplication). We will use B for this. Thus A is initially copied into B and B is set up for being accessed by the other processes as follows:

```
void bspsort(A,B,np) {
    bsp_begin(P);
    long p = bsp_nprocs();
    long s = bsp_pid();

    for(i=0;i<np;i++)
        B[i] = A[i];

    bsp_push_reg(B,np*sizeof(float));
    bsp_sync();

    // Your code goes here

    bsp_end();
}
```

- A) Write the code that counts the total number of elements across all processes that are smaller than each A[i]. The number for A[i] should be stored in pos[i] where pos is a local integer array of size np.
- B) Assume now that pos[i] contains the number of elements across all processes that are smaller than A[i]. Write the code that puts each element of A into its correct place of B on the appropriate process.

The syntax for bsp\_put is as follows:

```
bsp_put(pid, *source, *destination, offset_in_bytes, nbytes);
```

#### Problem 4, OpenMP

Consider the following parallel code for transposing an array using OpenMP. Assume that you run the code using two threads.

```
#pragma omp parallel for schedule(runtime)
for(int i=0;i<N;i++)
    for(int j=i+1;j<N;j++) {
        float temp = a[i][j]
        a[i][j] = a[j][i];
        a[j][i] = temp;
    }
```

- A. What is the (approximate) best speedup you can expect when using static scheduling?
- B. What is the (approximate) best speedup you can expect when using dynamic scheduling?