


# Parallel Programming Using Graphics Processing Units (GPUs)

# What is a GPU?

- Specialized hardware for processing graphics
- Available in almost all computers
- Simple and regular design
- Billions of transistors
- Can be integrated with the CPU or stand alone
- Performance: Up to 2,0 Gflops (2011)  
7.0 Tflops (2018)



## Main producers:

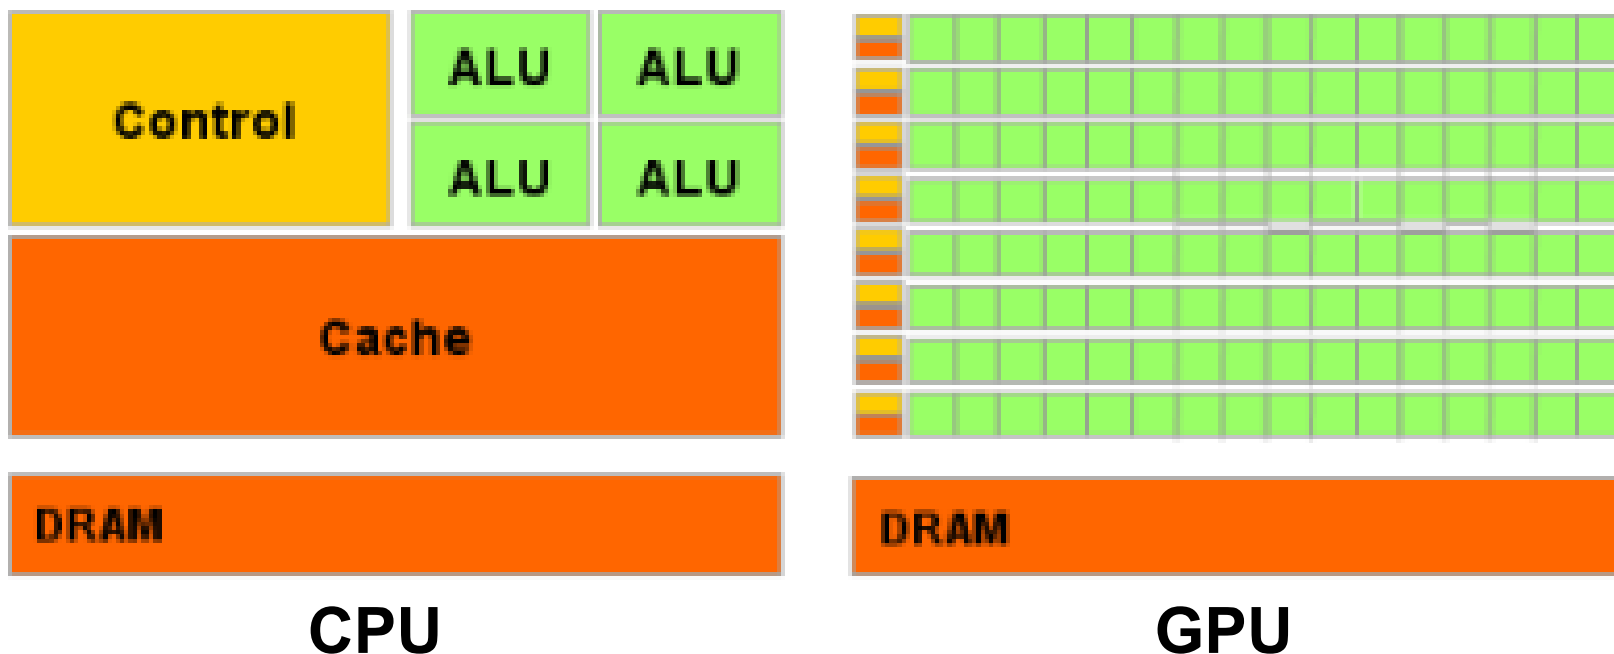
Nvidia, ATI (stand alone)

Intel (integrated)

ARM (mobile, in Trondheim!)



# Layout: CPU vs GPU



**DRAM:** memory

**Cache:** for data caching

**Control:** for flow control

**ALU** (Arithmetic-Logic Unit): for data processing

# Mode of Operation: CPU vs GPU

## **CPU**

- General purpose processing (Run Linux, file system, etc.)
- Typically, one thread on one CPU core
- Heavy weight threads: stop one thread to start another
- Threads on different cores can do different things (MIMD).

## **GPU**

- Specialized, only performs compute tasks
- Can run 1000+ threads simultaneously on one GPU processor
- Light weight threads: Hardware support for massive multi-threading.
- Works best when all the threads do the same thing (SIMD).

# Nvidia Tesla GPUs

- Designed for general computing
- A separate computing unit
- Up to 7 TFlops double precision
- Used in more than 50% of computers on Top500 list



# Programming the GPU

## **CUDA: Compute Unified Device Architecture (<http://www.nvidia.com/cuda>)**

- Designed by NVIDIA
- Supported on all modern NVIDIA GPUs (notebook GPUs, high-end GPUs, mobile devices)
- Gives high performance
- Forward compatible versions

## **OpenCL (<http://www.khronos.org/opencvl>)**

- Open standard, targeting NVIDIA, AMD/ATI GPUs, Cell, multicore x86, ..
- Gives portable code, can be run on all devices that support the model (CPU, GPU)
- Lower performance than CUDA
- Less mature

# CUDA – Compute Unified Device Architecture

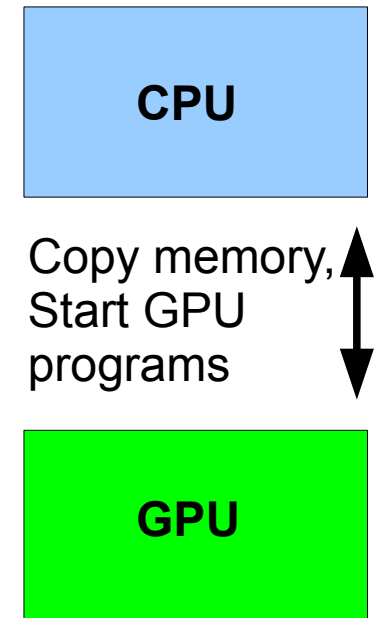
**Environment enabling programming of GPU h/w by a variety of common (non-graphics specific) high-level languages**

**Single source for CPU and GPU,**

- CPU starts up, use routines to transfer data to and from GPU.
- Language extensions used to control GPU routines (referred to as *kernels*)

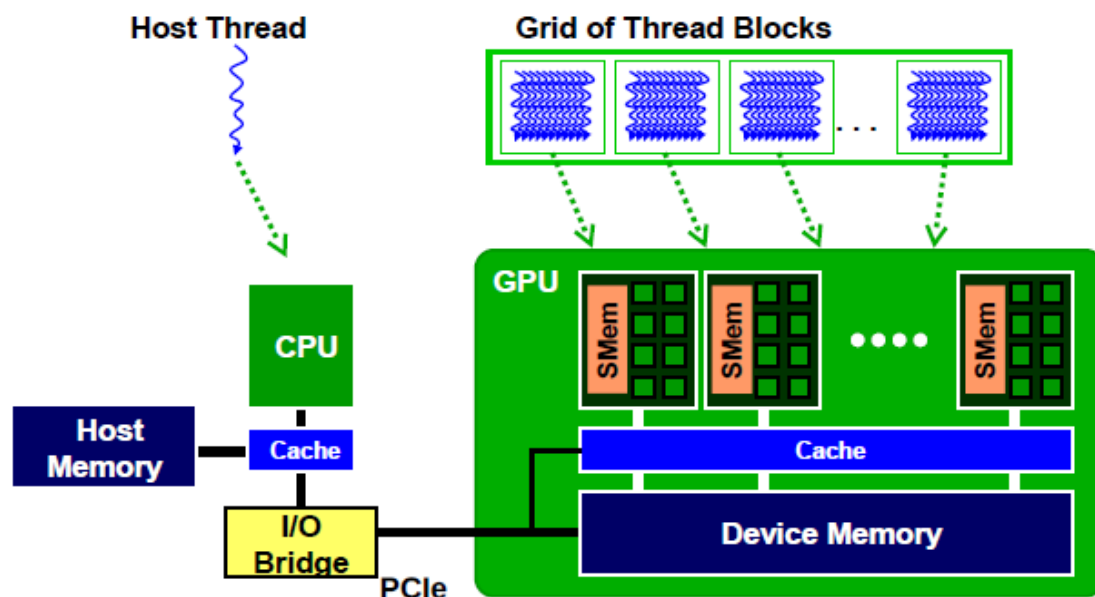
## **Languages**

- Host code C, C++, Fortran (PGI)
- GPU code C/C++ with extensions, Fortran with extensions
- `nvcc`: NVIDIA CUDA compiler
- PGI compiler for CUDA Fortran



# Program Outline


- (1) C program on CPU starts up
- (2) Performs initial (local) computation
- (3) Allocates memory on the GPU
- (4) Copies data from CPU memory to GPU memory
- (5) Specifies number of blocks and threads per block and starts GPU program (the *kernel*)
- (6) GPU program executes on GPU until completion
- (7) CPU copies result from GPU memory to CPU memory for post processing





# Organizing threads on the GPU

A thread block:




ID of a thread: `threadIdx.x`, `threadIdx.y`, `threadIdx.z`


```
myid = threadIdx.z*(blockDim.y*blockDim.x) + threadIdx.y*blockDim.x +
        threadIdx.x
```

If `blockDim.z = blockDim.y = 1` then

```
myid = threadIdx.x
```




# A grid of thread blocks



ID of a block:

`blockIdx.x`  
`blockIdx.y`  
`blockIdx.z`



If `threadDim.y = threadDim.z = 1` and  
`gridDim.y = gridDim.z = 1` then

`myid = blockIdx.x * blockDim.x +`  
`threadIdx.x`



# Kernel launch parameters

## Kernel launch parameters

- Block size: (x, y, z).  $x*y*z$  = Maximum of 768 threads total. (Hw dependent)
- Grid size: (x, y, z). Maximum of thousands of threads. (Hw dependent)

`dim3` is a Cuda data type for specifying dimension values.

Can take 1, 2, or 3 arguments:

```
dim3 blocks1D( 5          ); // Only x value, others = 1
dim3 blocks2D( 5, 5       ); // Only x and y, z = 1
dim3 blocks3D( 5, 5, 5 ); // x, y and z
```

Launching a kernel:

```
helloFromGPU<<<gridD, blockD>>>();
```

# The GPU Memory Model

## Local storage

- Each thread has own local storage
- Mostly registers, thus limited (KB)
- Data lifetime = thread lifetime
- Fast access

## Shared memory

- Each thread block has own shared memory
- Accessible only by threads within that block
- Data lifetime = block lifetime
- Small (KB), but fast

## Global (device) memory

- Accessible by all threads as well as host (CPU)
- Data lifetime = from allocation to deallocation
- Large (GB) but slower

# Basic C Extensions in CUDA

## Function modifiers

- `__global__`: to be called by the host but executed by the GPU.
- `__device__`: to be called and executed by the GPU only.

## Variable modifiers

- `__shared__`: variable in shared memory.
- `__device__`: variable in global memory
- Variables in kernels without modifiers are thread local

## Device functions

- `__syncthreads()`: sync of threads within a block.  
(similar to omp barrier)

# Example: Vector Addition

## Device Code

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    // Run grid of N/256 blocks of 256 threads each
    vecAdd<<< N/256, 256>>>>(d_A, d_B, d_C);
}
```

# Example: Vector Addition

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

Host Code

```
int main()
{
    // Run grid of N/256 blocks of 256 threads each
    vecAdd<<< N/256, 256>>>>(d_A, d_B, d_C);
}
```

Blocks    Threads

# Host Code for VecAdd (1)

```
// allocate and initialize host (CPU) memory  
float *h_A = ..., *h_B = ...; *h_C = ... (empty)
```

```
// allocate device (GPU) memory  
float *d_A, *d_B, *d_C;
```

```
cudaMalloc( (void**) &d_A, N * sizeof(float));  
cudaMalloc( (void**) &d_B, N * sizeof(float));  
cudaMalloc( (void**) &d_C, N * sizeof(float));
```

```
// copy host memory to device  
cudaMemcpy( d_A, h_A, N * sizeof(float), cudaMemcpyHostToDevice) );  
cudaMemcpy( d_B, h_B, N * sizeof(float), cudaMemcpyHostToDevice) );
```



# Host Code for VecAdd (2)

```
// execute grid of N/256 blocks of 256 threads each  
vecAdd<<<N/256, 256>>>(d_A, d_B, d_C);
```

```
// copy result back to host memory  
cudaMemcpy( h_C, d_C, N * sizeof(float), cudaMemcpyDeviceToHost) );
```

```
// do something with the result...
```

```
// free device (GPU) memory
```

```
cudaFree(d_A);  
cudaFree(d_B);  
cudaFree(d_C);
```

# VecAdd continued

What if N is not a multiple of 256?

Number of threads will always be a product of "block size" and "number of blocks"

N = 107, Number of blocks = 11, Number of threads per block = 10

```
int i = threadIdx.x + blockDim.x * blockIdx.x;
if (i < N)
    C[i] = A[i] + B[i];
```

What if "block size" \* "number of blocks" < N?

```
int i = threadIdx.x + blockDim.x * blockIdx.x;
while (i < N) {
    C[i] = A[i] + B[i];
    i += blockDim.x * gridDim.x;
}
```