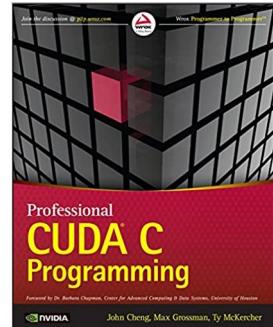
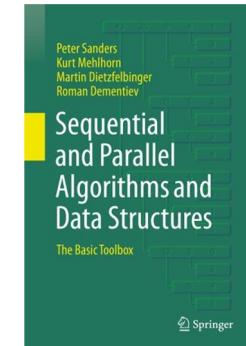


# INF236 – Parallel Programming

<b>Lecturer:</b>	Fredrik Manne	Fredrik.Manne@uib.no	Room 3145
<b>TA:</b>	FM & Kristian Sørdal	<a href="mailto:krisor99@gmail.com">krisor99@gmail.com</a>	
<b>Classes:</b>	Tuesday Thursday	08:15-10:00 10:15-12:00	Active room 3
<b>Exercises:</b>	Monday	08:15-10:00	Room 30202 (Rips)
<b>Grading:</b>	3 compulsory programming projects Final exam June 3rd, 09:00	50% of total grade 50% of total grade	
<b>Textbooks:</b>	<ul style="list-style-type: none"><li>- <i>Sequential and Parallel Data Structures and Algorithms, The Basic Toolbox</i> by Sanders, Mehlhorn, Dietzfelbinger, and Dementiev</li><li>- CUDA Programming : <i>Professional CUDA C Programming</i> by Cheng, John, et al.</li><li>- Handouts + online material</li></ul>		
<b>Information:</b>	<p><a href="http://mitt.uib.no">http://mitt.uib.no</a></p>		



# More information

Compulsory projects are individual work, unless otherwise specified.

This is strictly enforced!!!

Marcos Alonso Pardo  
Simon Andersen  
Thorgal Blanco  
Lars Bysheim  
Laxmi Dhital  
Østen Vestby Edvardsen  
Vegard Haugland  
Jonas Haukenes  
Thomas Svendsen Holden  
Oskar Aaltvedt Høgseth  
Magnus Haaland  
Gard Heine Kalland  
Sophie Koenig  
Felix Mariendal Kaasa  
Tarjei Landøy  
Fredrik Lauritzen  
Mariann Teigland Lepsøy

Moritz Nicolas Julian Makowski  
Olav Høysæther Opheim  
Jorge Pilan Saez Benito  
Sebastian Lillevik Refsnes  
Berenika Richterova  
Maya Alexandra Robbestad  
Jørgen Mykleby Storum  
Steffen Storøy  
Josef Sykora  
Markus Sandvær Sylta  
Kristian Sandvik Sørdal  
Lars Martin Taraldset  
Stian Aas Trohaug  
Carlos Vega de Alba  
Patrik Welters  
Yuhe Zhang

# Introduction

- Why we need powerful computers
- Trends in CPU development
- Why parallel computing is the only way to achieve performance
- Challenges when writing parallel algorithms

# But first some terminology..

*Big numbers:*  $10^3$  kilo,  $10^6$  mega,  $10^9$  giga,  $10^{12}$  tera,  $10^{15}$  peta,  $10^{18}$  exa

*Small numbers:*  $10^{-3}$  milli,  $10^{-6}$  micro,  $10^{-9}$  nano,  $10^{-12}$  pico,  $10^{-15}$  femto,  $10^{-18}$  atto

A computer running at 1 GHz:

- Performs 1 giga operations per second.
- Each operation is taking 1 nanosecond.
- Light travels 30cm in 1 nanosecond.

*Flops:* Number of 64-bit floating-point operations per second

One operation consists of one multiply and one add:

$$a = b + c * d$$

# So why do we need powerful computers?

# Scientific Applications

## Simulation: The third pillar of science

- Traditional scientific and engineering paradigm:

- 1) Do **theory** or paper design.
- 2) Perform **experiments** or build system.

- Limitations:

- Too difficult -- build large wind tunnels.
- Too expensive -- build a throw-away oil platform.
- Too slow -- wait for climate or galactic evolution.
- Too dangerous -- weapons, drug design, climate experimentation.

- Computational science paradigm:

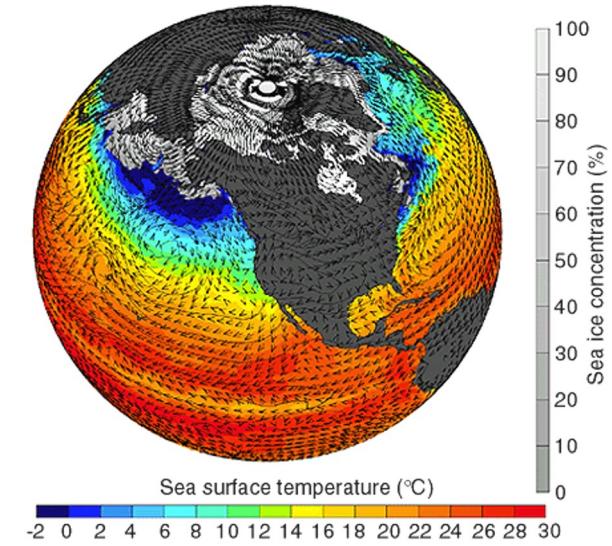
Use high performance computer systems to **simulate** the phenomenon

Based on known physical laws and efficient numerical methods.

# Weather and Climate Simulations

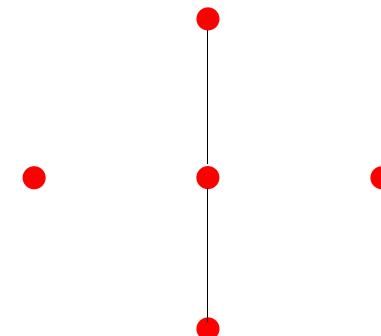
## Need model of interacting complex systems:

- Air: Wind, temperature, pressure, moisture, clouds
- Sea: Temperature, currents, salinity
- Land: Temperature, terrain
- Other: Sun, season, earth rotation, etc.



## Simulate continuous system:

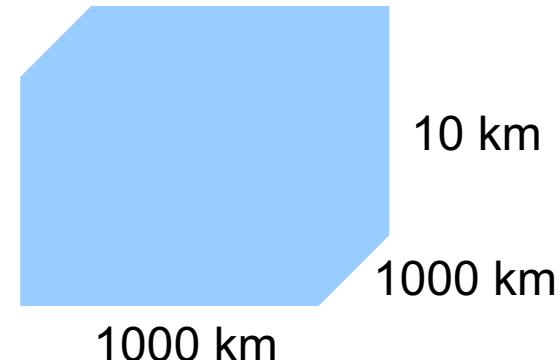
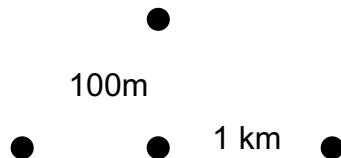
- Discretize system using a grid model
- What happens in one point is determined by what just happened in the surrounding points
- Apply initial conditions and run



# A Weather Model



- Simulate weather for 24 hours for a volume of 1000km x 1000km and of height 10km.
- Evenly spaced calculation points 1km apart and 100m above each other.
- Must make 100 calculations per point.
- Simulate for one day with 1.5 minute intervals  $\approx 1000$  time steps.



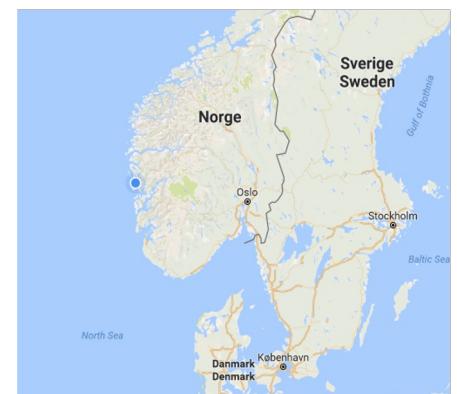
Total number of calculations =  $1000 * 1000 * 100 * 100 * 1000 = 10^{13}$  operations.

Computer running at 2GHz:

- One operation per clock cycle,
- $2 * 10^9$  operations per second
- Requires  $10^{13} / (2 * 10^9) = 5000$  seconds  $\approx 1.5$  hours

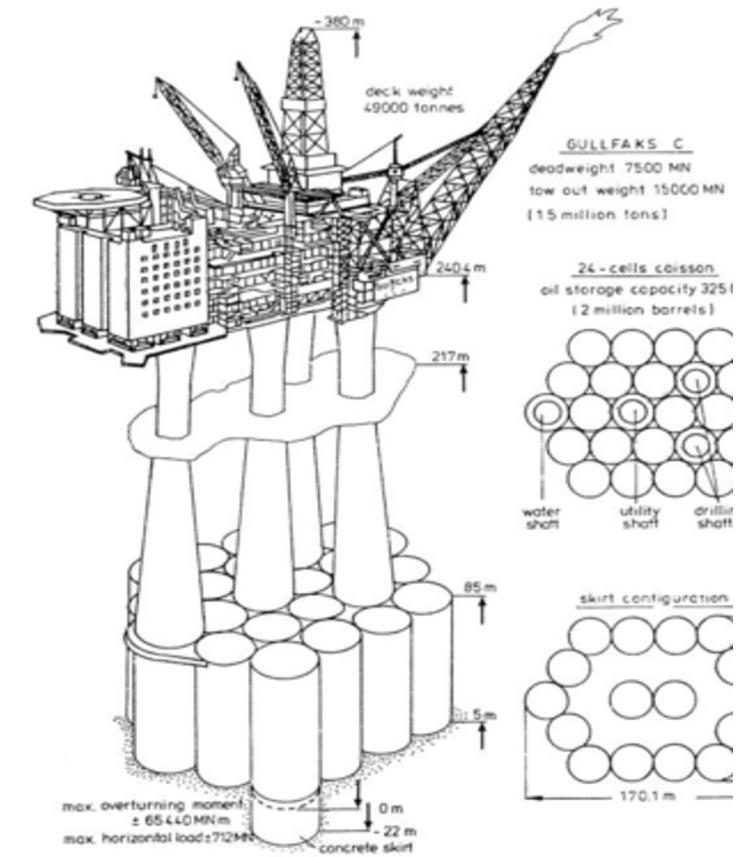
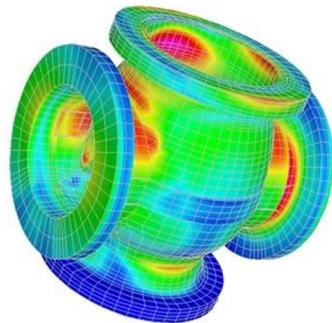
Refined mode: Put calculation points 100m apart in horizontal direction.

- Increases number of points to  $10000 * 10000 * 100 = 10^{10}$
- Requires  $100*1000*10^{10} / (2 * 10^9) = 0.5 * 10^6$  seconds  $\approx 139$  hours  $> 5$  days



# Engineering and Construction

- Discretize object depending on geometry.
- Apply various external forces and simulate.
- Each time step typically requires solution of system of partial differential equations (PDE).

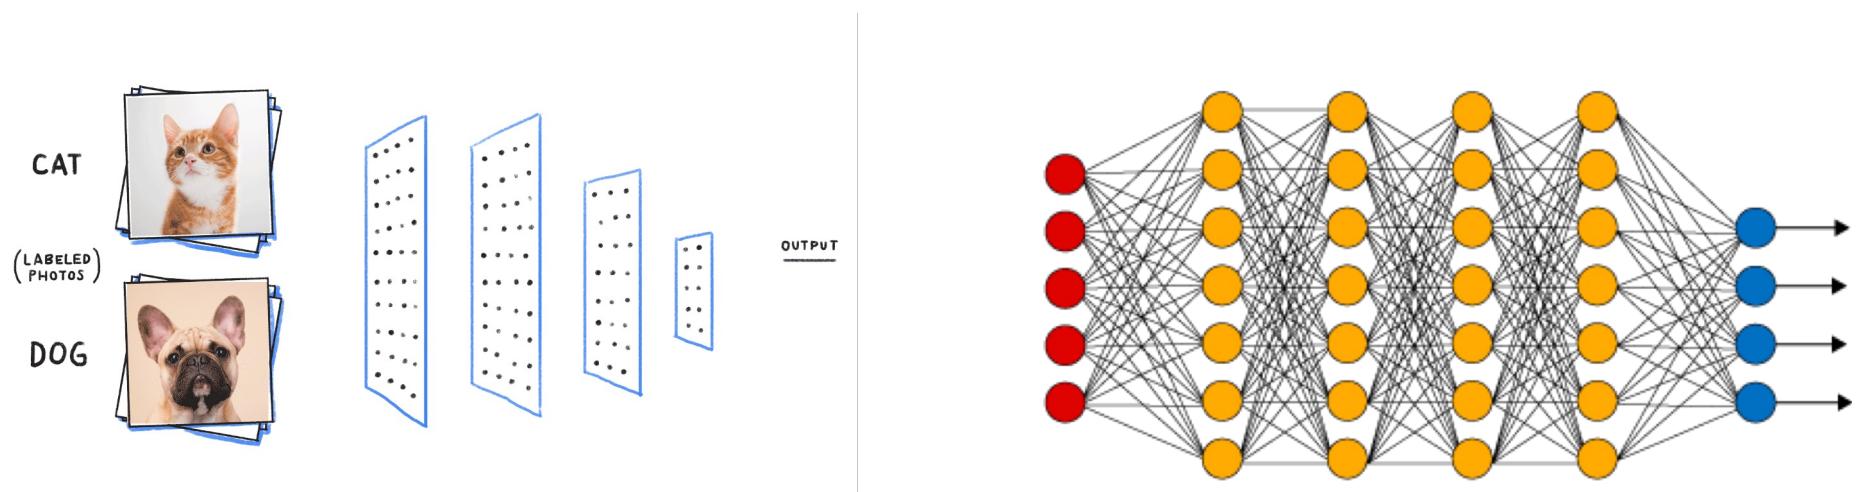


# Deep Learning

Machine learning method for classifying input

System is trained on large data sets, computationally demanding

Proposed around 1980, only became viable with the use of parallel processing



# 3D Computer Games

Maintains a 3D model and renders it on a 2D screen

Each pixel can be computed simultaneously

Driving force behind development of Graphical Processing Units (GPUs)

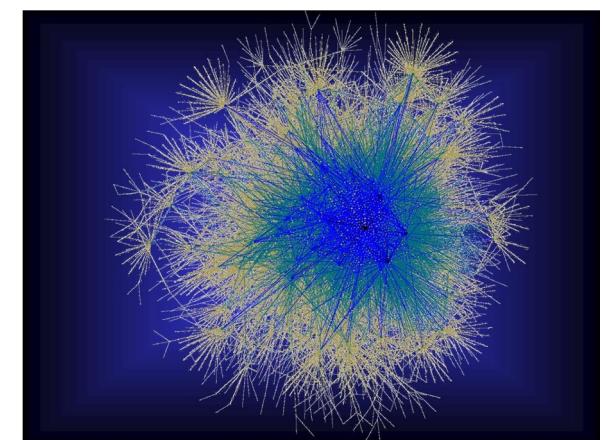
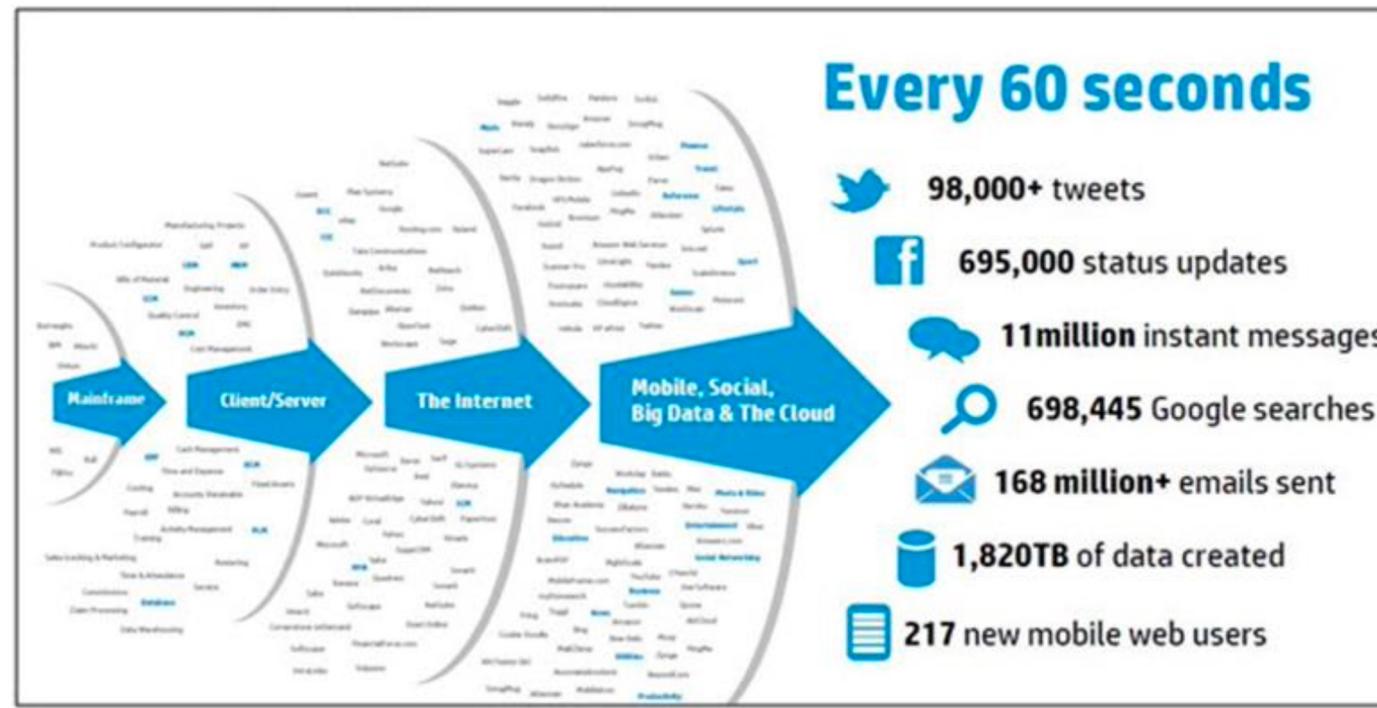


1992



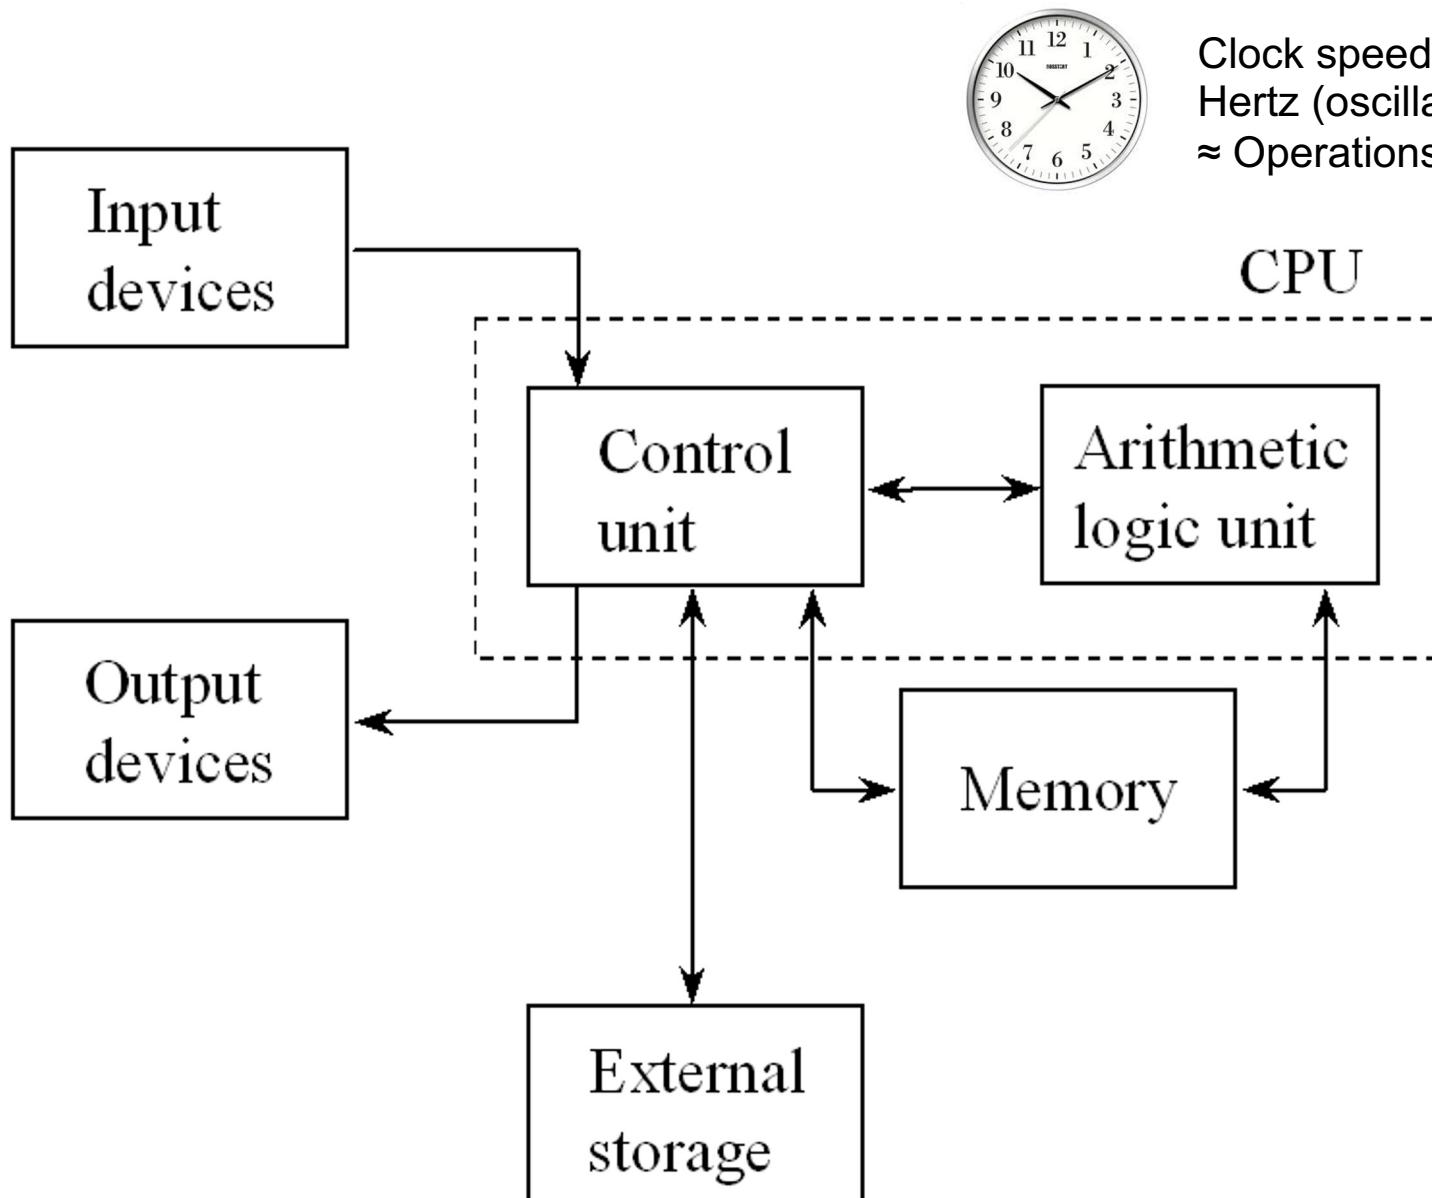
2021

# Data Explosion



# Trends in CPU development

# The Von Neumann Architecture



Clock speed is measured in Hertz (oscillations per second)  
≈ Operations per second

# Exponential Growth

Moore's law [1965]: The number of transistors that can be placed inexpensively on an integrated circuit doubles approximately every two years.



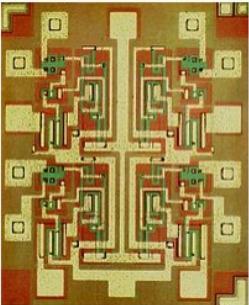
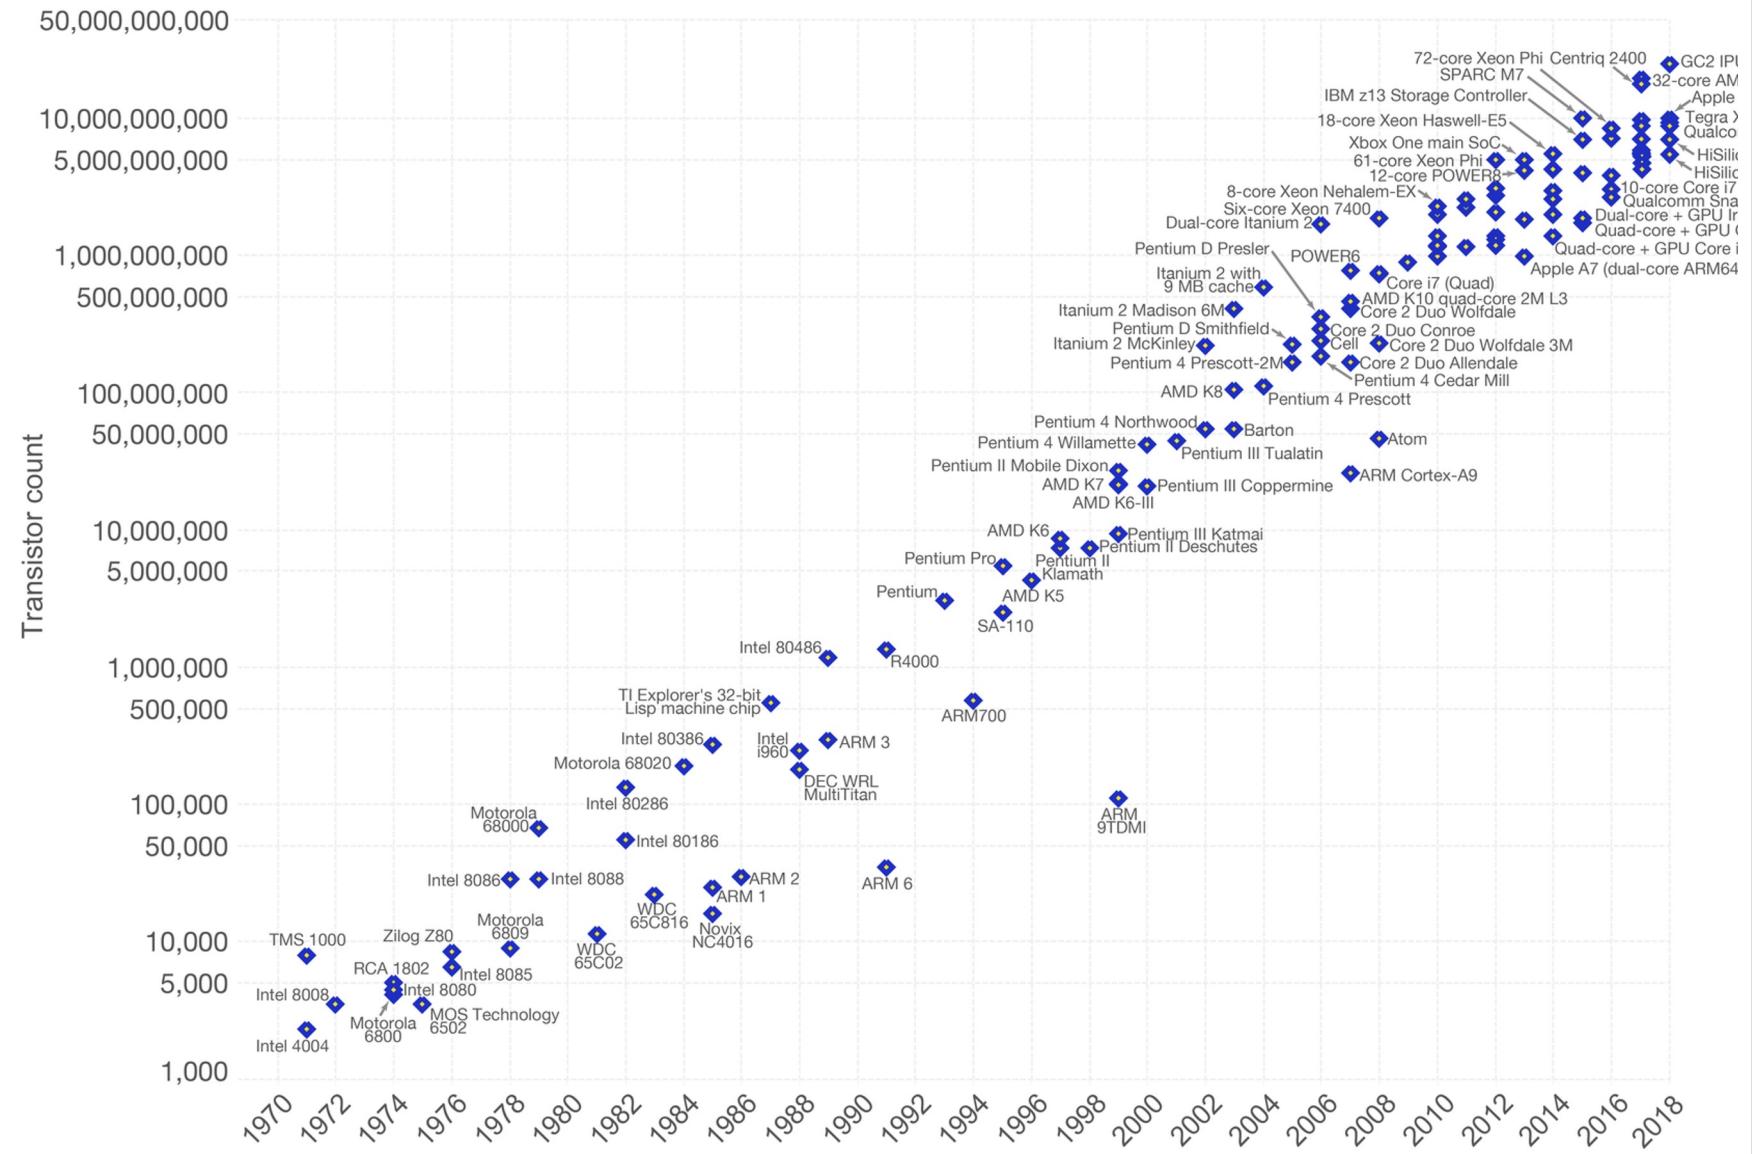
$$T = c * 2^{0.5k}$$

Applies to other areas as well:

- The clock rate of computers doubles every two years (used to be true).
- The amount of memory and the compute power of a computer, as well as the hard disk storage per dollar, grows exponentially.
- The capacity of optical fiber cables is doubled every nine months

## Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.

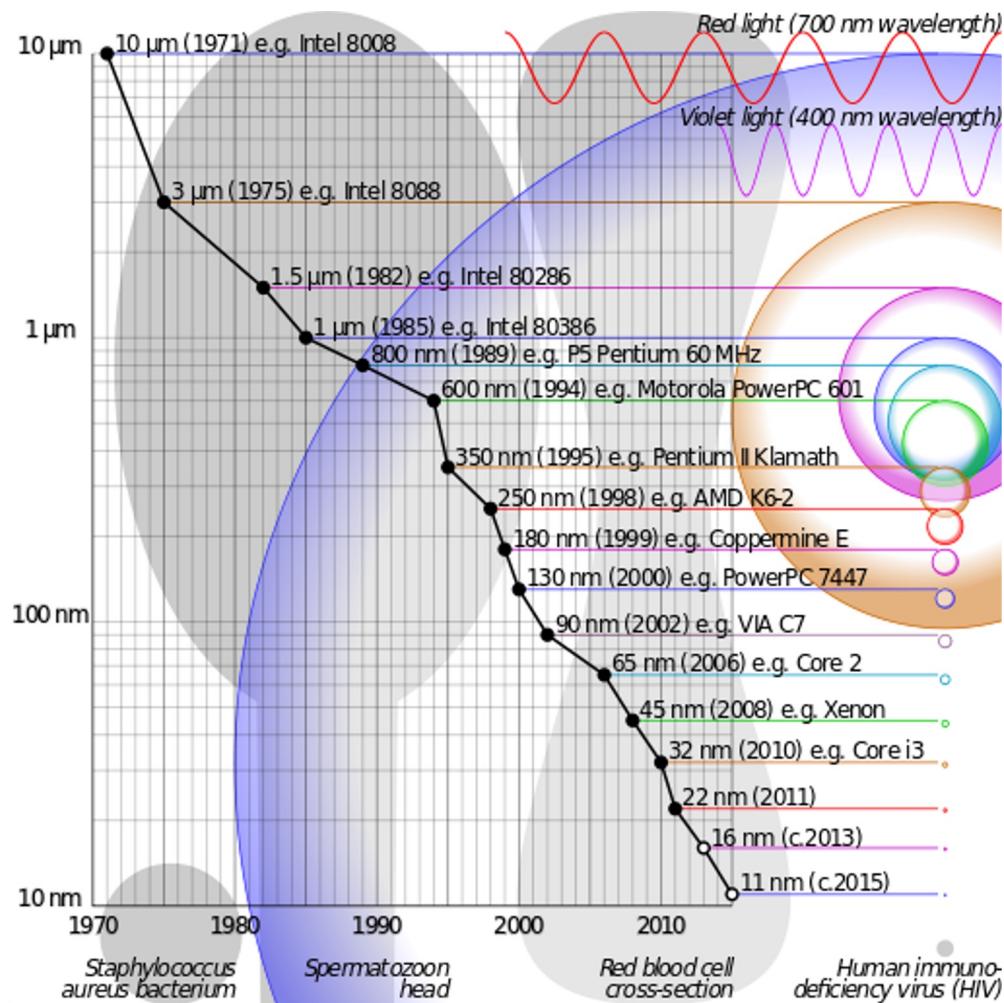


## MOSFET scaling (process nodes)

- 10  $\mu\text{m}$  – 1971
  - 6  $\mu\text{m}$  – 1974
  - 3  $\mu\text{m}$  – 1977
  - 1.5  $\mu\text{m}$  – 1981
  - 1  $\mu\text{m}$  – 1984
  - 300 nm – 1987
  - 600 nm – 1990
  - 350 nm – 1993
  - 250 nm – 1996
  - 180 nm – 1999
  - 130 nm – 2001
  - 90 nm – 2003
  - 65 nm – 2005
  - 45 nm – 2007
  - 32 nm – 2009
  - 22 nm – 2012
  - 14 nm – 2014
  - 10 nm – 2016
  - 7 nm – 2018
  - 5 nm – 2020
  - 3 nm** – 2022

## Future

# Transistor size



Nano meter (nm) =  $1 \times 10^{-9}$  meter

Helium atom = 0.1 nm

Hair width = 60 000 nm

A fingernail grows by 1 nm per second

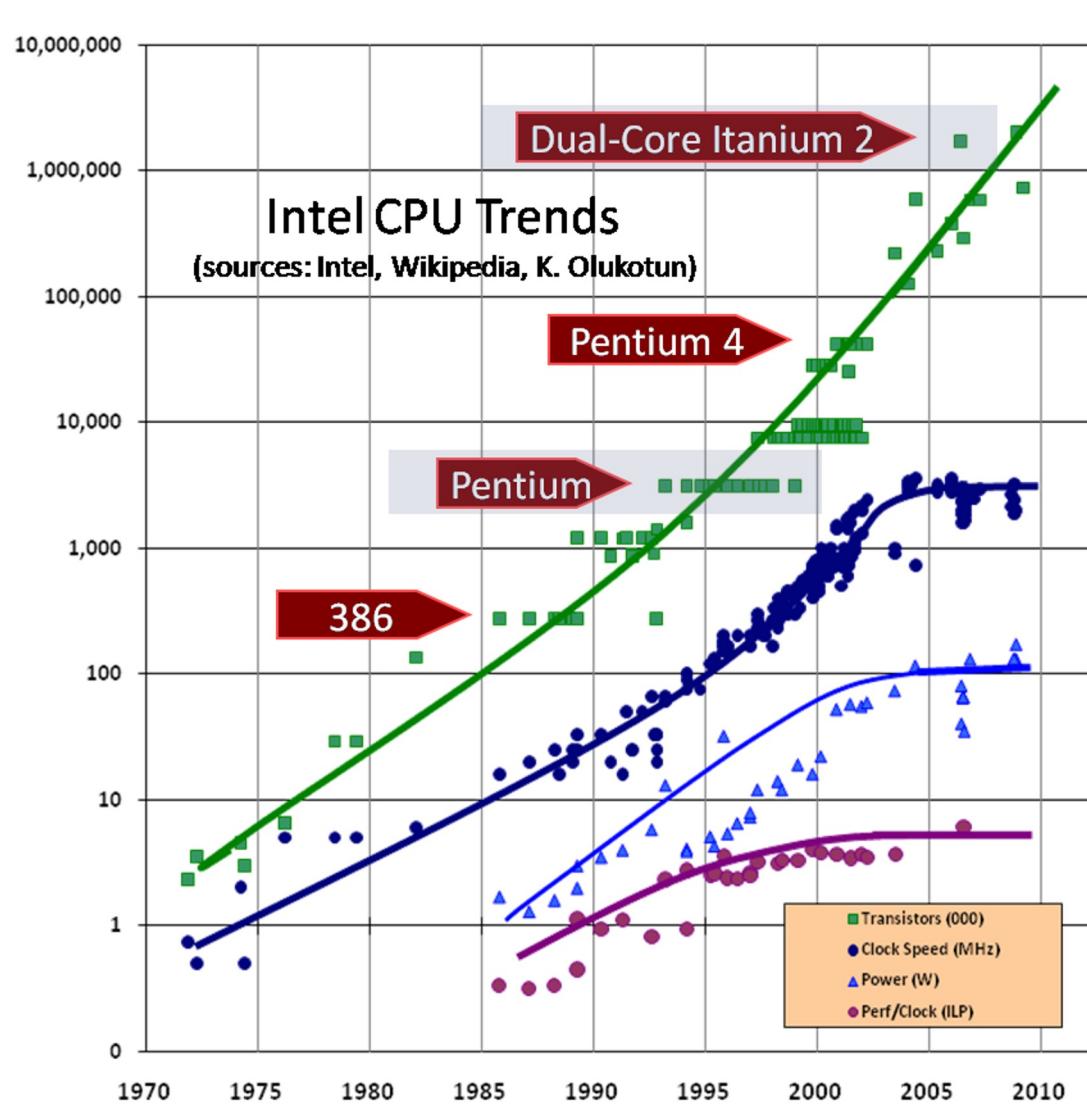
Machines used in this class:

brake.ii.uib.no: 32 nm technology.

lyng.ii.uib.no: 22 nm

birget.uib.no: 7nm

# Clock Frequency and Power



$$P = CU^2f$$

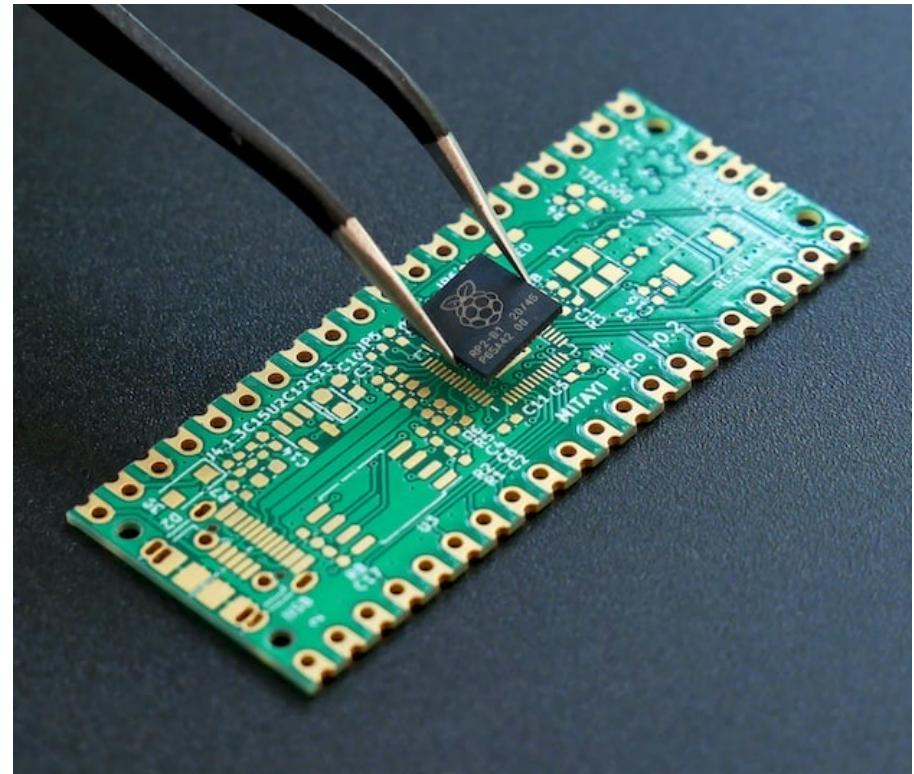
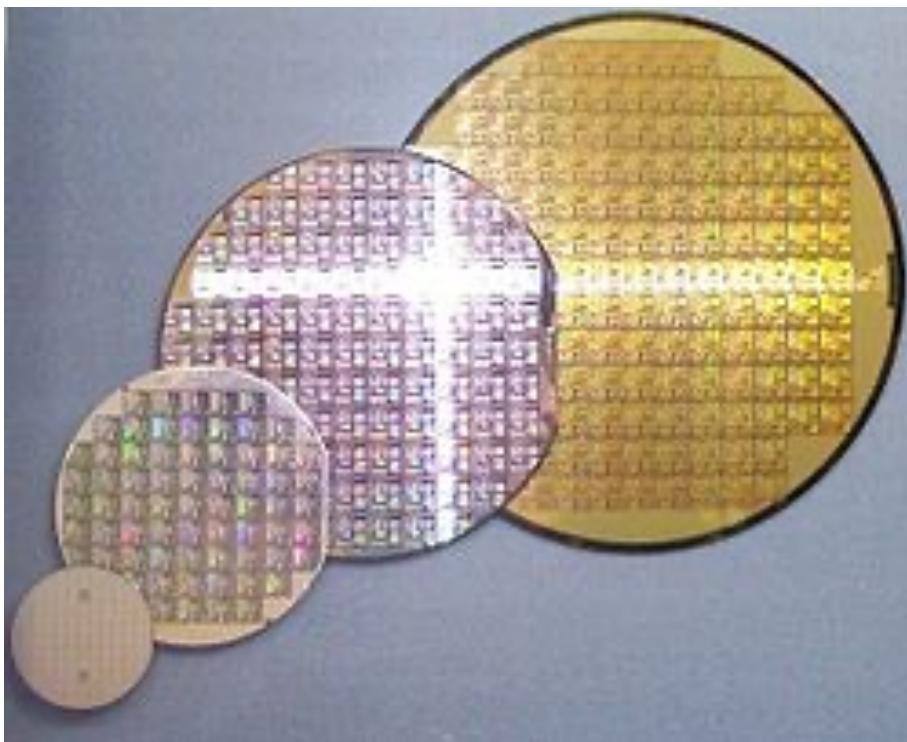
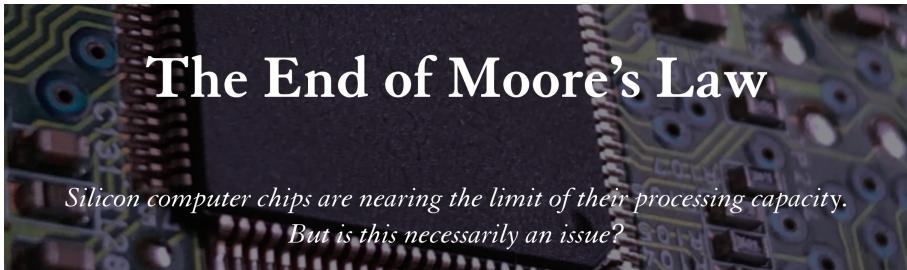
$P$ : Power  
 $U$ : Voltage  
 $f$ : frequency

Frequency grows with the voltage.  
Thus  $P \approx \Theta(f^3)$

Clock frequency has not increased (much) since 2000!

CPUs were getting too hot and using too much power.

# The Future



ANNALS OF TECHNOLOGY

# THE WORLD'S LARGEST COMPUTER CHIP

*In the race to accelerate A.I., the Silicon Valley company Cerebras has landed on an unusual strategy: go big.*

By Matthew Hutson

August 20, 2022

# Parallel Computing

# What to do?

What we have seen so far:

- Processor speed is not increasing (much)
- Transistor size is going down (at least for now)
- More space to fill with transistors: On the same area, count is going up as  $n^2$

What to do with the extra transistors?

- More memory
- More caches

But we still need more processing power!

Obvious solution:

Put multiple processors in the same computer!

# Parallel Computing is already here!



**Samsung Galaxy S21:**

2.2 GHz quad core processor (5nm)  
+ Mali- G78 GPU

**iMac:**

4.5GHz 4-18-cores + GPU

**Nvidia GeForce GTX 1660**

1536 cores (12 nm)

**Playstation 5:**

8 cores CPU (7 nm) + GPU

## Today

- PC's and phones: Mainly used for running multiple programs
- Graphic cards (GPU): Used to accelerate video and games (and ML)

## Future

- Clock rates are not increasing => Larger and more complex programs *must* use multiple processors
- Power goes down rapidly with decreasing frequency, remember  $P = \Theta(f^3)$   
→ Mobile devices can save battery by reducing clock rate while using more processors.

# Parallel Computing

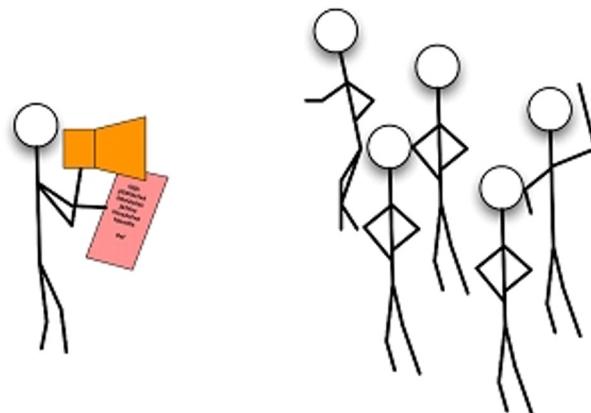
**Solve traditional computing problems using multiple interconnected processors**

## Why we do it:

- Get solution faster
- Gives access to more memory, can solve larger problems

## Issues:

- Harder to develop algorithms → “must find meaningful work for everyone”
- Harder to program → “do more than one thing at a time”
- Harder to debug → “execution is not deterministic”



NF236 - Parallel Co



# Other ways to exploit parallelism

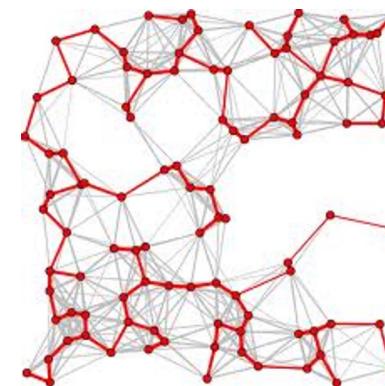
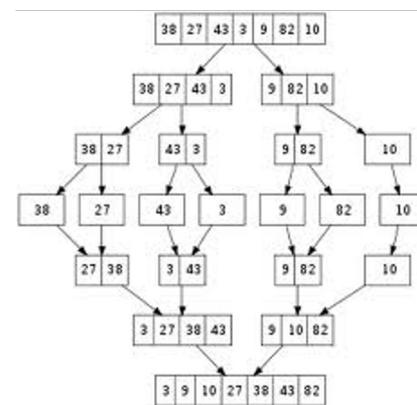
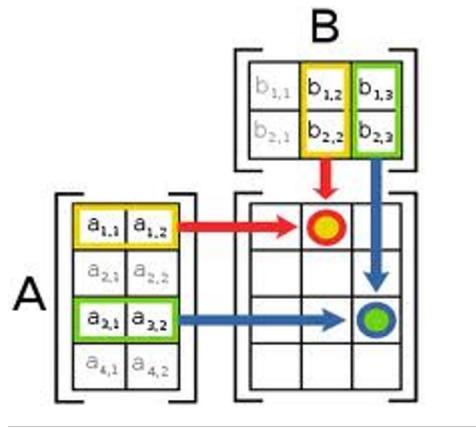
- Bit level parallelism
  - within floating point operations, etc.
- Instruction level parallelism (ILP)
  - multiple instructions execute per clock cycle
- Memory system parallelism
  - overlap of memory operations with computation
- OS parallelism
  - multiple jobs run in parallel

# Parallel Algorithms

We will develop and analyze parallel algorithms,

i.e., solve traditional algorithmic problems using parallel computers.

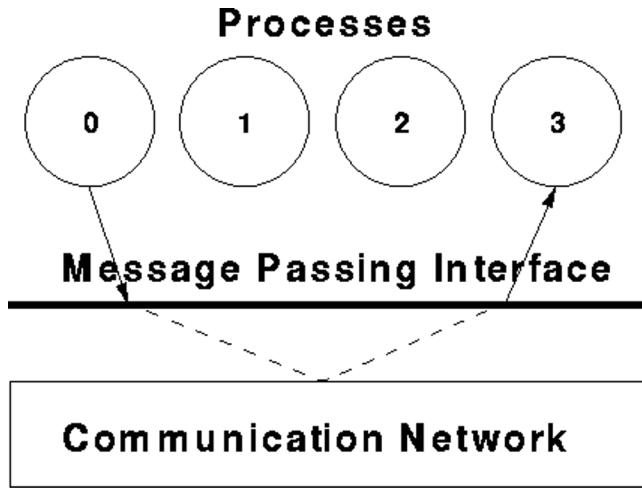
Main motivation comes from compute intensive applications:



We will not consider

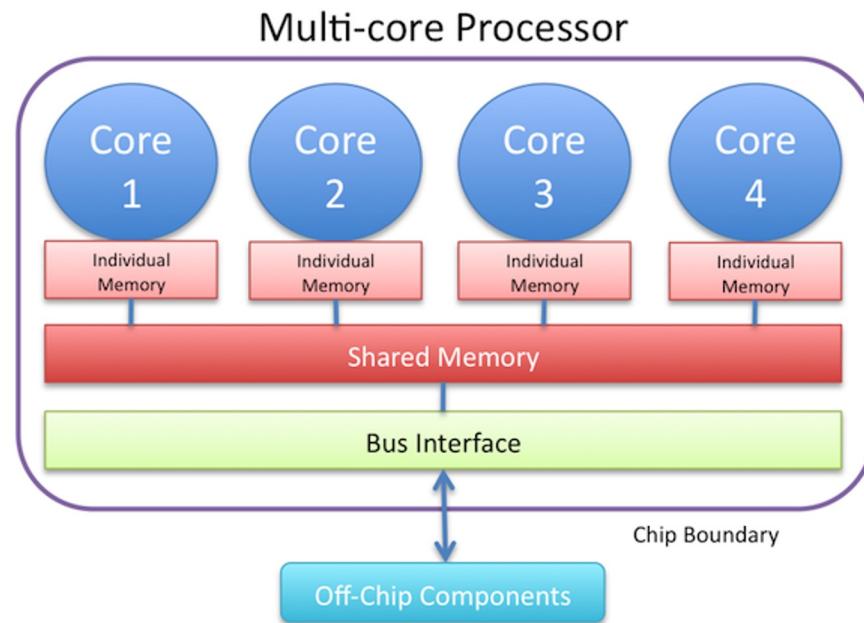
- Distributed computing
- Executing separate jobs on multiple processors (OS parallelism)

# Types of Parallel Computers

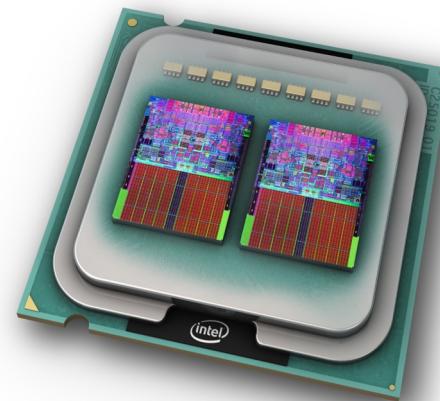


Communication Network

Distributed memory computers



Shared memory computers



# Top 500

Ranking of the 500 fastest computers in the world

Published twice a year

Based on solving  $Ax = b$  using Gaussian elimination with partial pivoting

Speed is measured in flops: Floating point operations per second

Current leader (2023):

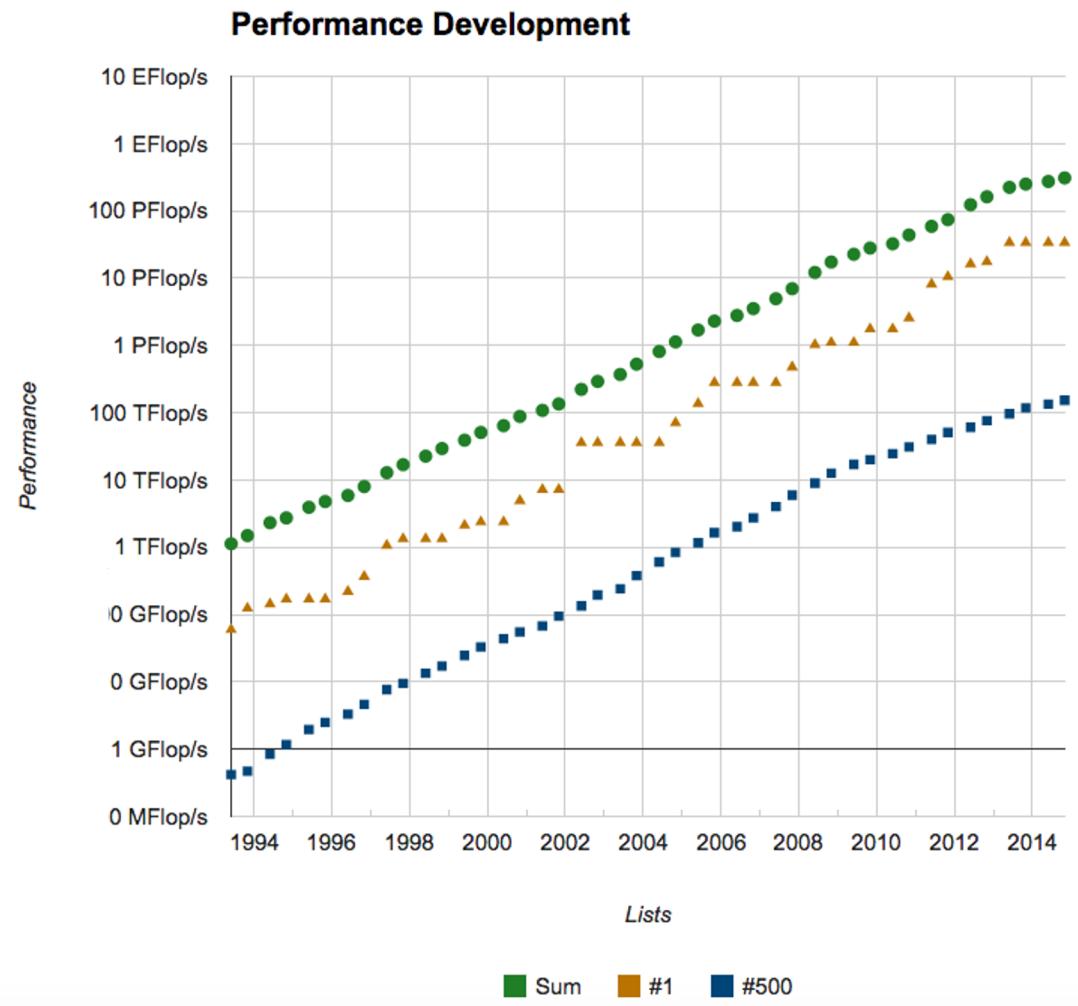
*Frontier* at Oak Ridge National Laboratories, USA  
8.7M Cores, 1102 petaflops (> 1 exaflop)

2022:

*Fugaku* at RIKEN Center for Computational Science, Japan  
7.6M Cores, 442 petaflops (peta =  $10^{15}$  0.44 exaflops)



Iphone 6



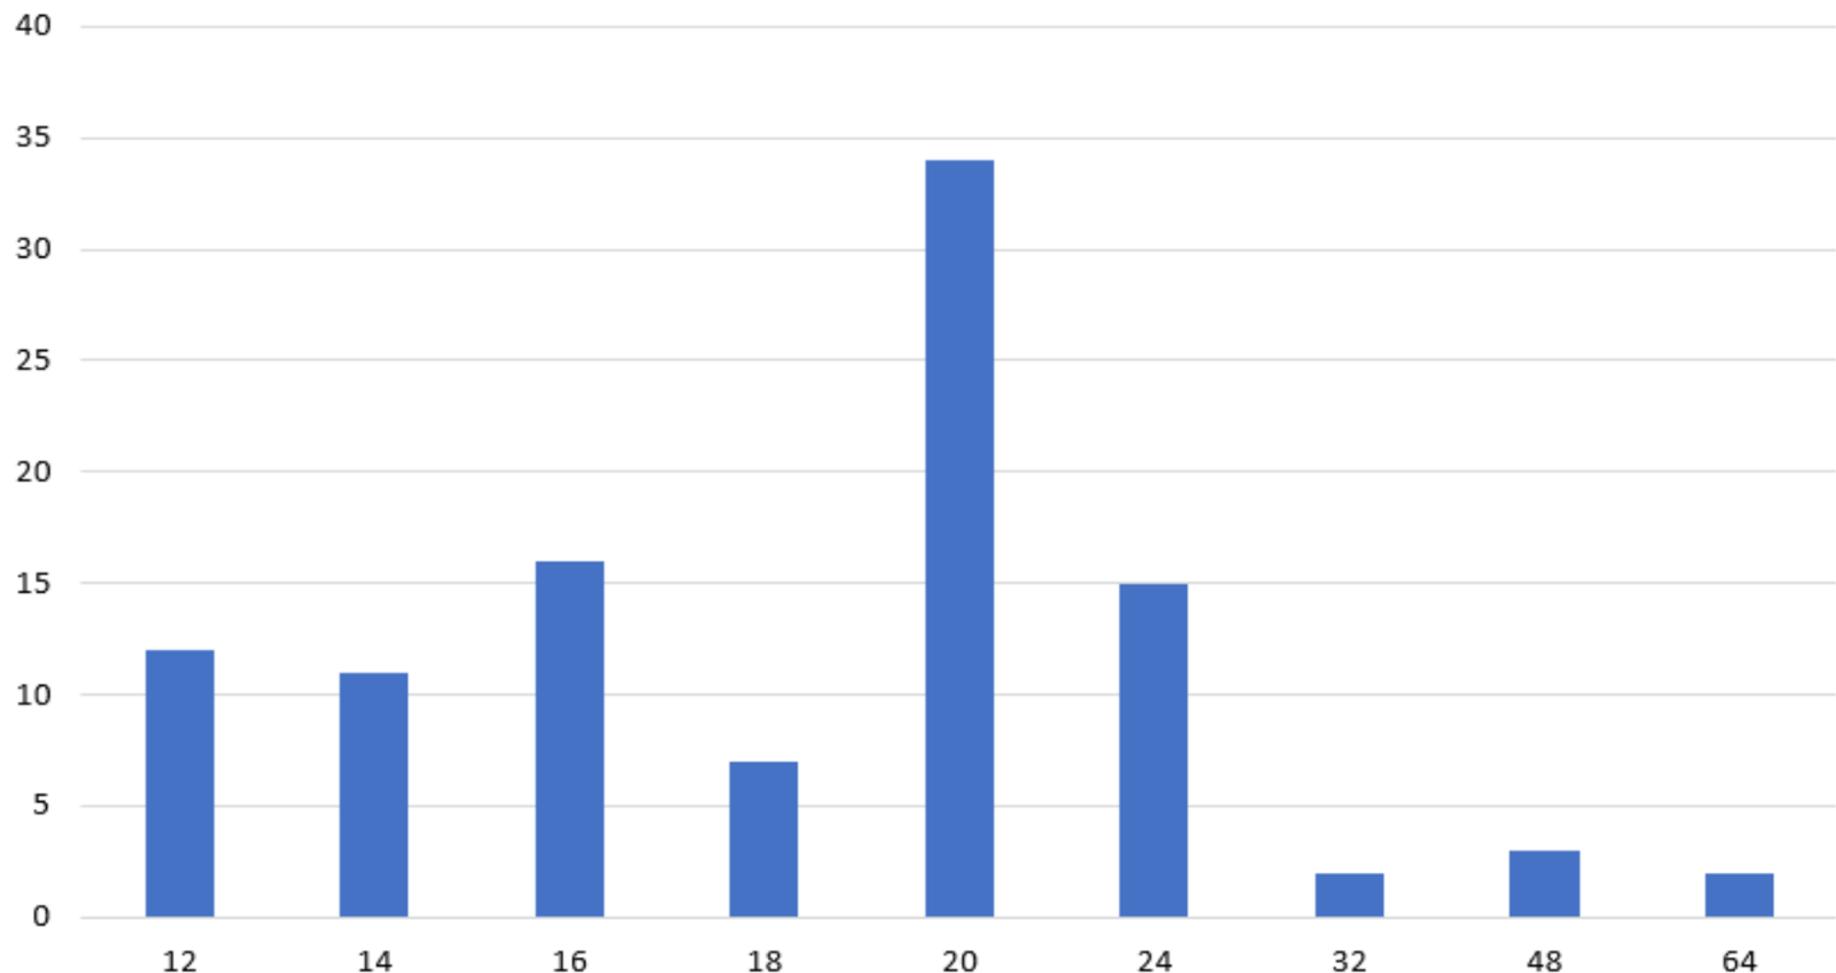
Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	<b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,730,112	1,102.00	1,685.65	21,100
2	<b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
3	<b>LUMI</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,220,288	309.10	428.70	6,016
4	<b>Leonardo</b> - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail NVIDIA HDR100 Infiniband, Atos EuroHPC/CINECA Italy	1,463,616	174.70	255.75	5,610
5	<b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148.60	200.79	10,096
6	<b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox	1,572,480	94.64	125.71	7,438

# November 2019 Top500 New Systems By CPU Cores Per Socket

**November 2019 Top500 Supercomputers**



New Systems CPU Cores Per Socket



November 2019 Top500 New Systems By CPU Cores Per Socket

# Hexagon

- Most powerful computer in Norway 2008 - 2012
- Highest ranked Norwegian computer ever (48, June 2008)



	Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
Nov 08	66	University of Bergen Norway	Cray XT4 QuadCore 2.3 GHz Cray Inc.	5,550	40.6	51.1	274
	82	University of Tromso Norway	<b>Stallo</b> - Cluster Platform 3000 BL460c, Xeon 53xx 2.66GHz, GigEthernet Hewlett-Packard	5,632	31.9	59.9	
Nov 09	121	University of Bergen Norway	Cray XT4 QuadCore 2.3 GHz Cray Inc.	5,550	40.6	51.1	274
	193	University of Tromso Norway	<b>Stallo</b> - Cluster Platform 3000 BL460c, Xeon 53xx 2.66GHz, GigEthernet Hewlett-Packard	5,632	31.9	59.9	
Nov 10	255	University of Bergen Norway	Cray XT4 QuadCore 2.3 GHz Cray Inc.	5,550	40.6	51.1	274
	363	StatOil Norway	Cluster Platform 3000 BL460c G6, Xeon X5570 2.93 GHz, GigE Hewlett-Packard	5,664	35.3	66.4	
	458	University of Tromso Norway	<b>Stallo</b> - Cluster Platform 3000 BL460c, Xeon 53xx 2.66GHz, GigEthernet Hewlett-Packard	5,632	31.9	59.9	
Nov 12	56	Norwegian University of Science and Technology Norway	SGI ICE X, Xeon E5-2670 8C 2.600GHz, Infiniband FDR SGI	22,048	396.7	458.6	537
	134	University of Oslo Norway	<b>Abel</b> - MEGWARE MiriQuid, Xeon E5-2670 8C 2.600GHz, Infiniband FDR Megware	10,080	178.6	209.7	227
	157	University of Bergen Norway	<b>Hexagon</b> - Cray XE6m-200, Opteron 6276 16C 2.300GHz, Cray Gemini interconnect Cray Inc.	22,272	160.1	204.9	342.3

Nov 13

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
82	Norwegian University of Science and Technology Norway	SGI ICE X, Xeon E5-2670 8C 2.600GHz, Infiniband FDR SGI	22,048	396.7	458.6	537
240	University of Oslo Norway	Abel - MEGWARE MiriQuid, Xeon E5-2670 8C 2.600GHz, Infiniband FDR Megware	10,080	178.6	209.7	227
281	University of Bergen Norway	Hexagon - Cray XE6m-200, Opteron 6276 16C 2.300GHz, Cray Gemini interconnect Cray Inc.	22,272	160.1	204.9	342.3

Jun 14

RANK	SITE	SYSTEM	RMAX		RPEAK		POWER
			CORES	(TFLOP/S)	(TFLOP/S)	(KW)	
98	Norwegian University of Science and Technology Norway	SGI ICE X, Xeon E5-2670 8C 2.600GHz, Infiniband FDR SGI	22,048	396.7	458.6	537	
300	University of Oslo Norway	Abel - MEGWARE MiriQuid, Xeon E5-2670 8C 2.600GHz, Infiniband FDR MEGWARE	10,080	178.6	209.7	227	
382	University of Bergen Norway	Hexagon - Cray XE6m-200, Opteron 6276 16C 2.300GHz, Cray Gemini interconnect Cray Inc.	22,272	160.1	204.9	342.3	

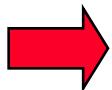
Nov 14

RANK	SITE	SYSTEM	RMAX		RPEAK		POWER
			CORES	(TFLOP/S)	(TFLOP/S)	(KW)	
127	Norwegian University of Science and Technology Norway	SGI ICE X, Xeon E5-2670 8C 2.600GHz, Infiniband FDR SGI	22,048	396.7	458.6	537	
368	University of Oslo Norway	Abel - MEGWARE MiriQuid, Xeon E5-2670 8C 2.600GHz, Infiniband FDR MEGWARE	10,080	178.6	209.7	227	
462	University of Bergen Norway	Hexagon - Cray XE6m-200, Opteron 6276 16C 2.300GHz, Cray Gemini interconnect Cray Inc.	22,272	160.1	204.9	342.3	

# Developing Parallel Algorithms

# Challenges when writing parallel programs

- Finding enough parallelism (Amdahl's Law)
- Granularity
- Locality
- Load balance
- Coordination and synchronization

 All these challenges make parallel programming even harder than sequential programming.

# Until next time...

- Check out [www.top500.org](http://www.top500.org)
- Check out [www.green500.org](http://www.green500.org)
- Check out [www.graph500.org](http://www.graph500.org)
- Intel's tick-tock cycle of development
- Read about Gordon Moore and John Von Neumann
- Read about Cerebras