

# Arrays, queues and linked lists

”Problems that are relatively easy to parallelize”

# Trivially Parallelizable Problems

## Characteristics:

- Independent computations
- Each computation of (approximately) the same size
- Easy to partition the computations

## What to keep in mind:

- Partition computation on the highest possible level
- Is it worth it? Overhead vs. speedup
- Shared memory: Careful with false sharing (cache effects)
- Distributed memory: Need enough computation to offset cost of communication

# Matrix Operations

$C = A + B$  where  $A$ ,  $B$ , and  $C$  are  $n \times n$  matrices

```
for i = 1,n  
  for j = 1,n  
    c(i,j) = a(i,j) + b(i,j)
```

$$t_s = n^2, t_p = n^2/p + t_{\text{comm}}$$

## OpenMP:

- Parallelize outermost loop
- Little risk of false sharing
- Different access pattern might change how matrices are stored

## Message passing on distributed memory computer

Speedup depends on where  $A$  and  $B$  are stored initially and where  $C$  should be stored.

$A$ ,  $B$ , and  $C$  on process 0:

- $t_p$  will most likely be higher than  $t_s$

$A$ ,  $B$ , and  $C$  distributed:

- $t_{\text{comm}} = 0$

# Dynamic Scheduling

## Setting

- Unknown job sizes
- Different processor speeds
- Unknown number of jobs

## Shared memory

- Use guided or dynamic load balancing
- Can also use `task` directive

## Distributed memory

- Master thread sends out tasks upon request
- Each worker thread receives tasks, computes, and then sends back answer.
- The job size must be determined empirically

# Example: Recursive Functions

**Recursive functions**  $s(n) = f(s(n-1), s(n-2) \dots s(0))$

**Fibonacci sequence:**  $s(n) = s(n-1) + s(n-2)$ , where  $s(0) = 0$ ,  $s(1) = 1$

**2-dimensional functions:**

Takes a two-dimensional point as input  $(a,b)$

**Complex numbers:**

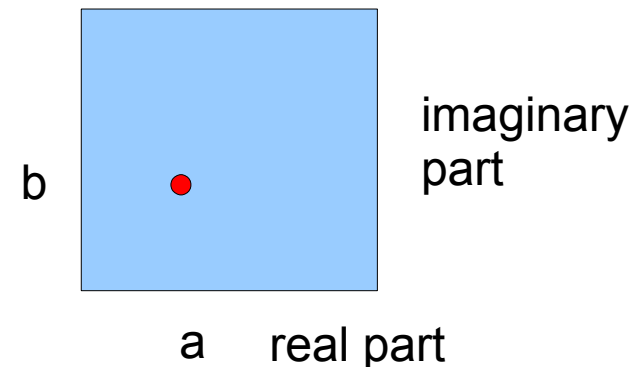
$c = a + bi$  where  $i = \sqrt{-1}$

$a$  is the *real* part,  $b$  is the *imaginary* part

$s(n) = (s(n-1))^2 + c$  where  $s(0) = 0$

**Question:** For which values of  $c$  will  $s(n)$  diverge and for which will the value stay bounded?

Complex coordinate system



# The Mandelbrot Set

$$s(n) = (s(n-1))^2 + c \text{ where } c = a + bi \text{ and } s(0) = 0$$

If  $s(n)$  does not diverge then  $c$  is in the Mandelbrot set.

**Fact:** If  $S(n) = x + yi$  then  $s(n)$  will diverge iff  $|s(n)| = \sqrt{(x^2 + y^2)} \geq 2$ .

$$\begin{aligned}(s(n))^2 &= (x+yi)^2 = x^2 + 2xyi + y^2i^2 \\ &= x^2 - y^2 + 2xyi\end{aligned}$$

$$\begin{aligned}s(n+1) &= (s(n))^2 + a + bi \\ &= (x^2 - y^2 + a) + (2xy + b)i \\ &\quad \text{Real part} \quad \text{Imaginary part}\end{aligned}$$

```
while (i < limit) and (x2 + y2) < 4
    {t = (x2 - y2 + a); y = (2xy + b); x = t; i++;}
```

```
if i == limit
    then (a,b) is in the M.S.
```

# Plotting the Mandelbrot Set

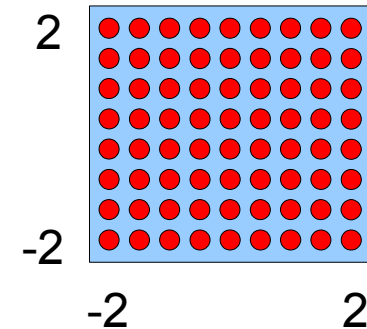
**Area of computation is continuous**

→ Need to discretize.

**Speed of convergence varies**

→ Number of computations depends on coordinates

Use colors to indicate speed of divergence



## Going parallel

Computation in each point is independent

→ Easy to parallelize with shared memory

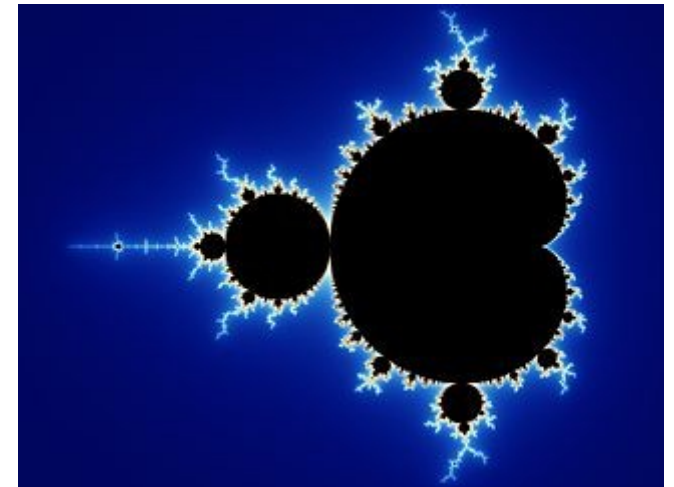
## **Distributed memory**

→ Partitioning of the work only requires coordinate information

→  $O(1)$  data out.

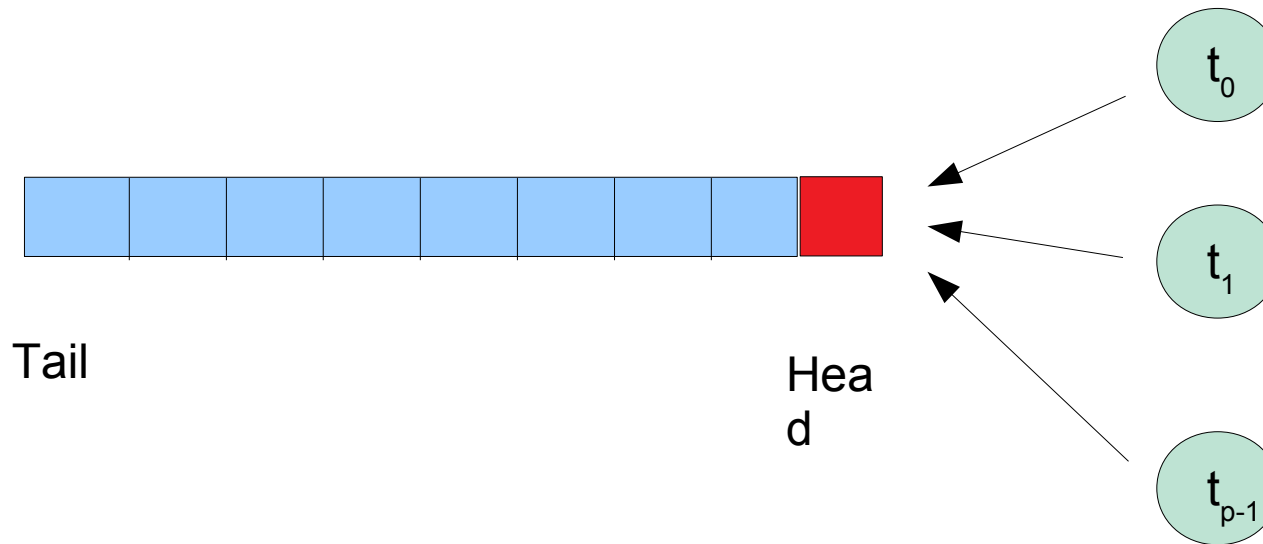
$O((n^2 * \text{max it.})/p)$  computations

$O(n^2/p)$  data in.



# Queues

Parallel access to a common queue can be difficult



Must serialize access to Head (or Tail)

A thread needs to both read and write to Head pointer

## Solutions:

- Use a software lock (`omp_lock`)
- Use hardware dependent methods (Compare-and-swap)
- Could try with one queue for each thread?

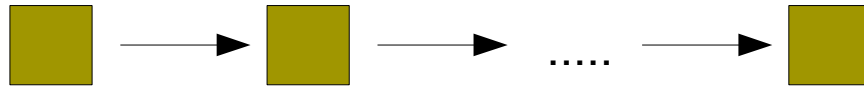


Still, introduces sequential part...



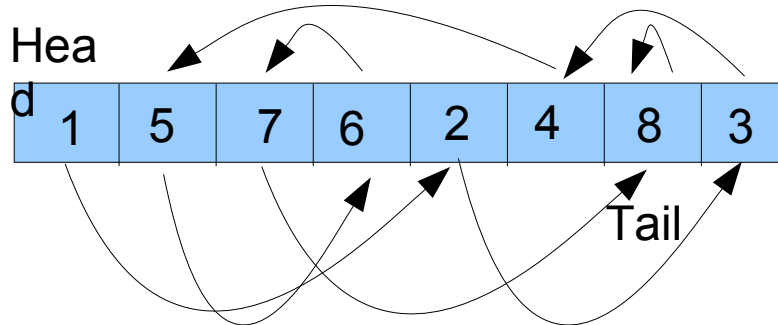
# Lists

**General case: Can only access Head (and maybe Tail)**



Head  
Tail  
Can only be accessed sequentially :-)

**But what if the list is stored in an array?**

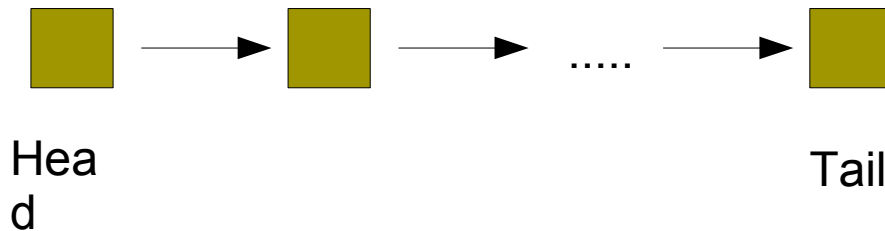


**Could we sort the elements in parallel?**

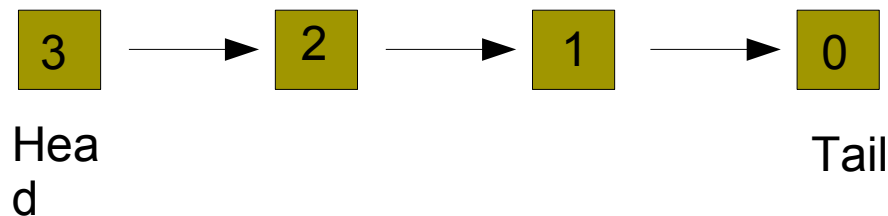
Might be more suitable for parallel processing.

# List Ranking Problem

**Need to compute where to put each element in the final array**



For each element count the number of ensuing elements.  
Can easily compute final position and move elements in parallel.



(See book or later lecture for actual algorithm)

# Monte Carlo Simulations

Monte Carlo methods tend to follow a particular pattern:

- Define a domain of possible inputs.
- Generate inputs randomly from a probability distribution over the domain.
- Perform a deterministic computation on the inputs.
- Aggregate the results.

Usefull for search problems where finding an exact value is too costly

# A Simple Example

**Given**  $n$  numbers  $a_1, a_2, \dots, a_n$ .

**Problem:** Find a number larger than the median.

Traditional algorithm: Must look at at least  $n/2 + 1$  values.

Thus takes time  $O(n)$

**Monte Carlo Method:**

Choose  $k$  values at random

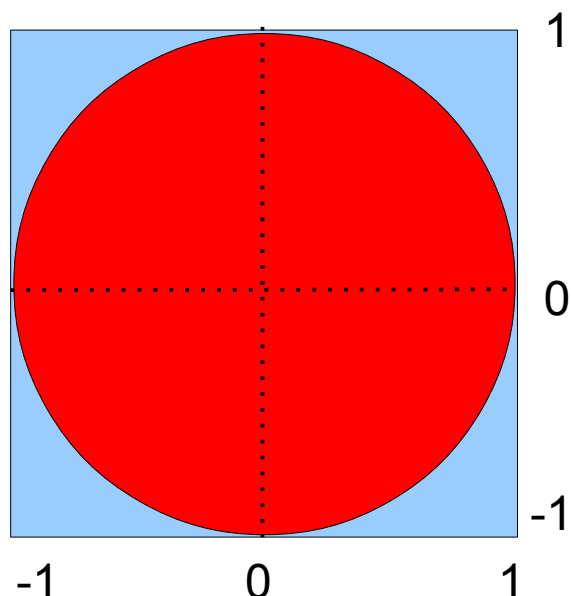
Pick max value out of the  $k$  and return. Time  $O(k)$

**How good is it?**

$$\begin{aligned} P(\text{max value in upper half}) &= 1 - P(\text{max value in lower half}) \\ &= 1 - P(\text{all } k \text{ values in lower half}) \\ &= 1 - (0.5)^k \end{aligned}$$

$k = 10$  gives  $P(\text{success}) > 0.999$ ,  $k = 20$  gives  $P(\text{success}) > 0.9999999$

# Calculating $\pi$



$$\text{Area of circle} = \pi \times r^2 = \pi$$

$$\text{Area of square} = 4$$

$$\text{Area of circle} / \text{Area of square} = \pi / 4$$

Choose random point in the square

Register number of points inside of the circle.

$$(x,y): \quad x^2 + y^2 > 1 \rightarrow \text{Outside of circle}$$

$$x^2 + y^2 \leq 1 \rightarrow \text{Inside of circle}$$

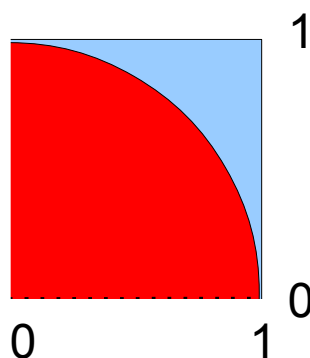
$$\text{Number inside} / \text{Total number} \rightarrow \pi/4$$

**Easy to parallelize!!**

Could also use only one quadrant

$$\text{Area of circle} = \pi/4$$

$$\text{Area of square} = 1$$



# Generating random numbers

## Pseudorandom numbers:

- Recursion:  $\mathbf{x}_{i+1} = \mathbf{f}(\mathbf{x}_i)$ ,  $\mathbf{x}_0 = \text{seed}$ 
  - $\mathbf{f}$  is chosen such that " $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$  appear to be random"
  - Guessing  $\mathbf{x}_{n+1}$  should be "difficult" when  $\mathbf{x}_0$  is unknown
- Reproducible sequence:
  - Use a **deterministic** seed
- Unpredictable sequence:
  - Use a hash function value and current **time** as seed
- `int random(void); // Returns the next pseudorandom number`
- `void srand(unsigned int seed); // Sets the seed`

## Linear congruential generators (Knuth 1981):

- Recursion:  $\mathbf{x}_{i+1} = (\mathbf{a}\mathbf{x}_i + \mathbf{c}) \% \mathbf{m}$ 
  - Parameters  $\mathbf{a}$ ,  $\mathbf{c}$  and  $\mathbf{m}$  should be chosen to give the sequence a long **period**
  - After one period, the sequence will repeat itself.

# Generating random numbers

## Pseudorandom numbers in parallel programs:

- Centralized: A dedicated thread generates all pseudorandom numbers
  - May involve many requests to the generating thread
  - Not very suitable for Monte Carlo simulation
- Distributed: Each thread generates the numbers as they are needed
  - Sequences must be distinct
  - Seeds must be distinct
  - Can use a hash function of (time and) process rank

`srandom()`, `random()`: Sequential methods for seeding and generating random numbers, all numbers are generated from the same sequence.

`srandom_r()`, `random_r()`: Thread safe, generates different sequences on different threads. Slightly more complicated to use.