

Sorting

"The Good, the Bad and the Ugly"

Today: The Bad!

Sorting Algorithms

A number of $\Theta(n^2)$ algorithms:

- Selection sort
- Bubble sort
- Insertion sort

Several $O(n \log n)$ algorithms:

- Quicksort
- Mergesort
- Heapsort

Other sorting algorithms

- Counting sort
- Radixsort
- Bucketsort

Selection sort

Given array $A[0:n-1]$ of unsorted numbers
Order elements of A in increasing order

General idea:

Find the smallest element, swap with the first element
Repeat $n-2$ times

For $i = 0$ to $n-2$

$\text{smallest} = A[i]$

$\text{position} = i$

 For $j=i+1$ to n

 If $A[j] < \text{smallest}$

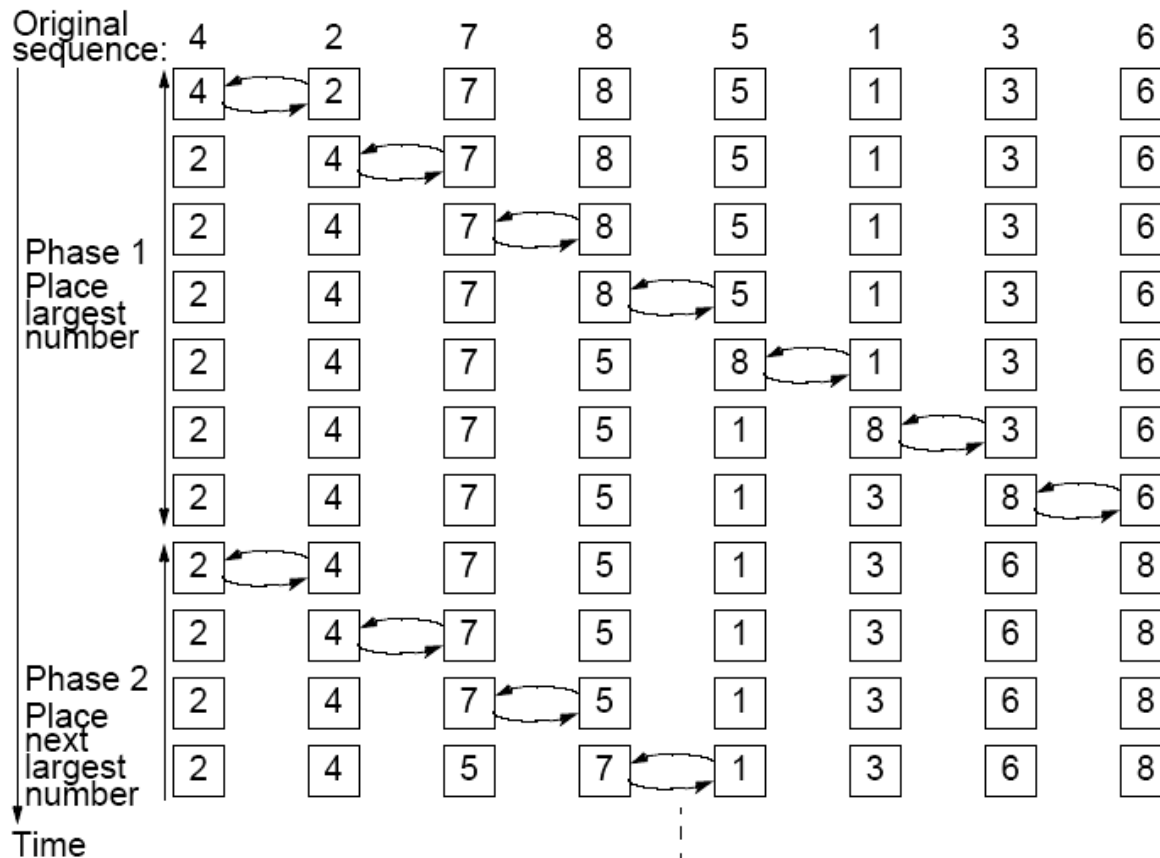
$\text{smallest} = A[j]$

$\text{position} = j$

 Swap $A[i]$ and $A[\text{position}]$

Bubble sort

For $i = n-1$ to 1 // Bubble $i+1$ 'st largest element to index i
 For $j = 1$ to i
 compare-swap a_j and a_{j+1}



Time:

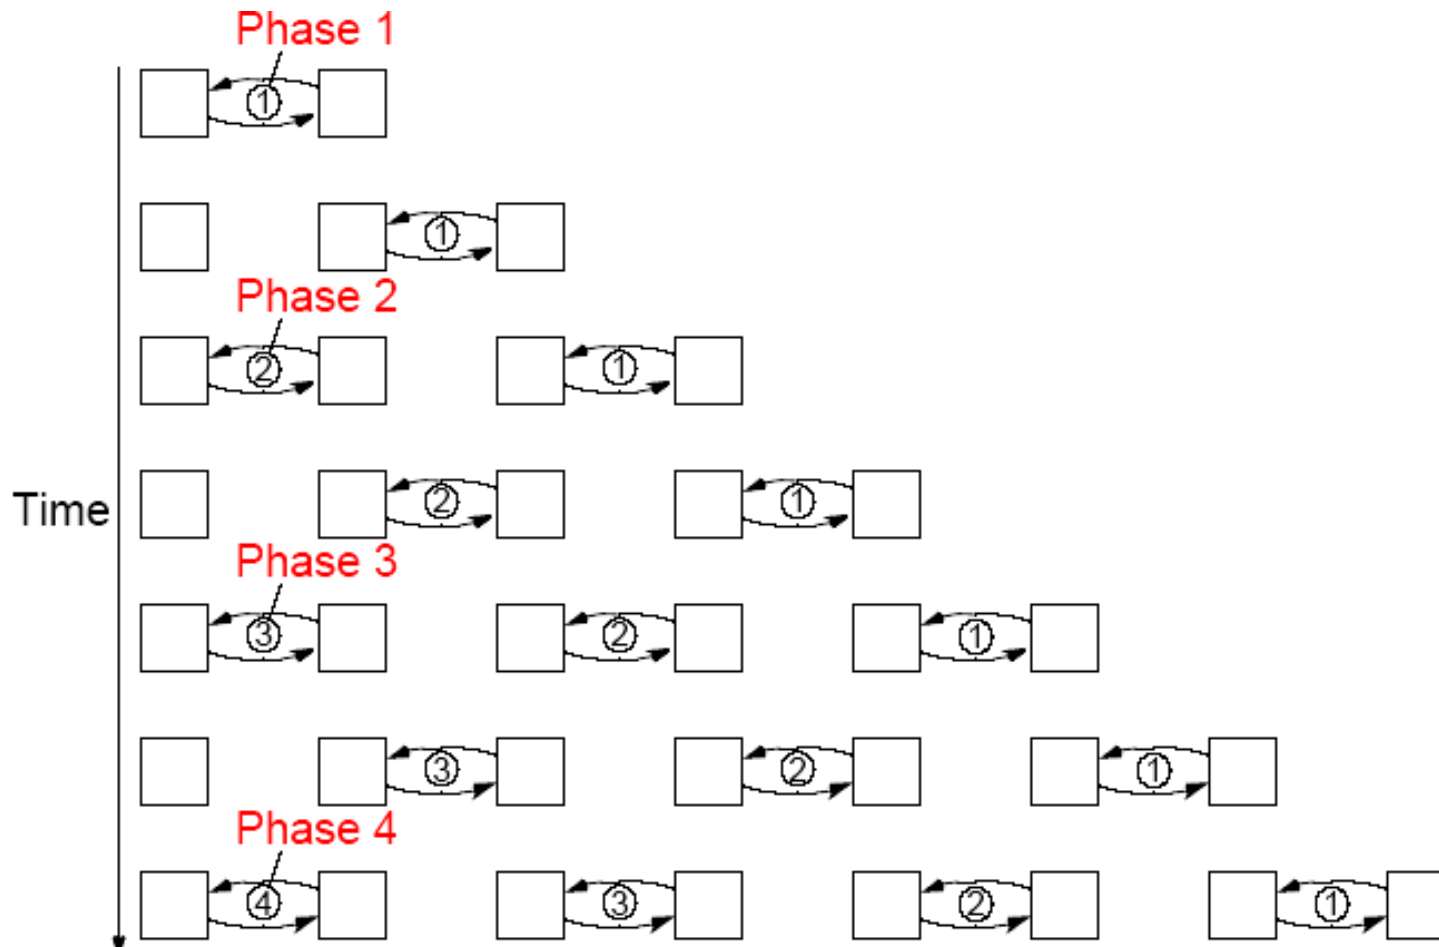
$$\sum_{i=1}^{n-1} i = n(n-1)/2 = \Theta(n^2)$$

Note:

With one thread per element the second stage could start earlier.

Parallel Bubble Sort

Phase i starts as soon as possible as long as it does not overtake iteration $i+1$.



With $p = n$ the algorithm takes time $\Theta(n)$.

Perfect speedup compared to sequential bubblesort,
but only $\log(n)$ compared to mergesort.


Odd-Even Transposition Sort

"Use all the threads in each round."

For step = 1, n

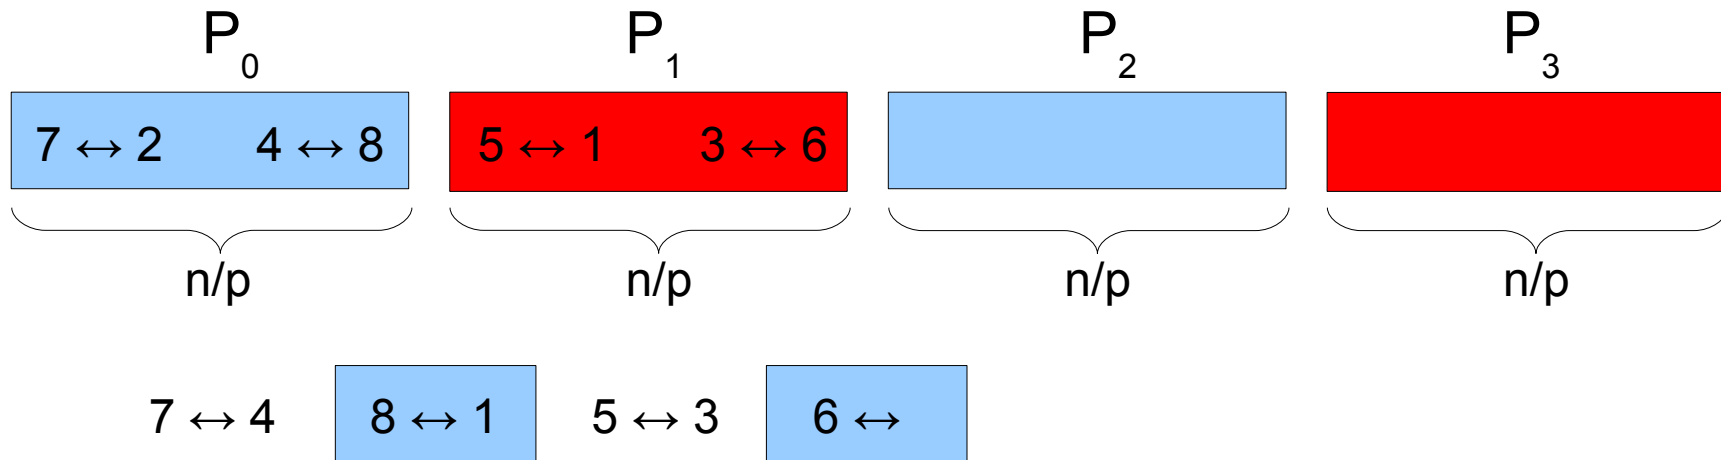
Even step: Even p_i talks to p_{i+1} , Odd p_i talks to p_{i-1}

Odd step: Even p_i talks to p_{i-1} , Odd p_i talks to p_{i+1}



Practical Implementation

Since $p \ll n$ each thread must handle n/p elements in each iteration.




n iterations, $O(n/p)$ work in each iteration


$$\rightarrow T_p = O(n^2/p)$$

Practical Implementation

Since $p \ll n$ each thread must handle n/p elements in each iteration.



Could move anywhere from 1 to n/p elements between processes



Practical Implementation

Better strategy: Pre-sort lists of $n/(2p)$ elements each
Perform $2p$ steps of block merge, each taking time $\Theta(n/p)$.

$$T_p = \underbrace{O(n/p \log(n/p))}_{\text{Presort}} + \underbrace{\Theta(n)}_{2p \cdot n/p}$$

Note:


- Can stop when there is no data movement
- Could perform initial long range movements
- Might make subsequent algorithm faster

Divide and Conquer

From sequential algorithms:

Quicksort

Work is in
partitioning



Mergesort

Work is in putting
things together

Tasks on each level are typically independent

In parallel:

Divide until parts are of size n/p , and then use a sequential algorithm
Can also divide in more than two parts on each level.


Parallel Mergesort

Sequential:

```
for(i=1; i<n; i*=2)
  merge two and two
  adjacent lists of length i
```

- $\log n$ levels
 - merge takes $O(n)$ time on each level
- $\rightarrow t_{\text{seq}} = O(n \log(n))$

Parallel:



Level Length



Parallel Running Time

1. stage. Every processor sorts n/p elements, $T1_p = O(n/p \log(n/p))$

2. stage. Length of lists doubles for every level,

- starts with length n/p , ends with length n
- Length = $n/p, 2n/p, 4n/p, \dots, n$
- Total of $\log p$ steps

Thread 0 dominates running time of second stage

$$\begin{aligned} T2_p &= 2^1 n/p + 2^2 n/p + 2^3 n/p \dots + 2^{\log(p)} n/p \quad (\text{note } 2^{\log(p)} = p) \\ &= n/p * (2^1 + 2^2 + 2^3 \dots + 2^{\log(p)}) = n/p * (2p-1) = 2n - n/p \end{aligned}$$

$$T_p = T1_p + T2_p = 2n + n/p (\log(n/p) - 1)$$

$$S = T_{\text{seq}} / T_p \approx n \log(n) / (2n + n/p \log(n/p)) < 1/2 \log(n)$$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] \leq a[1] \leq \dots \leq a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$

$[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{\}, \{\}, \{\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{\}, \{9\}, \{\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{\}, \{9\}, \{15\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$

$[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3\}, \{9\}, \{15\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9\}, \{15\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9, 6\}, \{15\}, \{\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9, 6\}, \{15\}, \{21\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9, 6\}, \{15\}, \{21, 18\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:
 $\{3, 0\}, \{9, 6\}, \{15, 12\}, \{21, 18\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9, 6\}, \{15, 12\}, \{21, 18\}$

Step 1: Sort buckets:

$\{0, 3\}, \{6, 9\}, \{12, 15\}, \{18, 21\}$

Bucket Sort

Input: n integers $a[0..n-1]$ (evenly distributed) on $[0, b-1]$

Output: $a[0] < a[1] < \dots < a[n-1]$

Example ($n=8, b=24$): $a = \{9, 15, 3, 0, 6, 21, 18, 12\}$

Algorithm:

Divide interval $[0, b-1]$ into m buckets of equal width:

Example ($m=4$): Bucket width = $b/m = 24/4=6$
 $[0,6), [6,12), [12,18), [18,24)$

Step 0: Put $a[0..n-1]$ in buckets:

$\{3, 0\}, \{9, 6\}, \{15, 12\}, \{21, 18\}$


Step 1: Sort buckets:

$\{0, 3\}, \{6, 9\}, \{12, 15\}, \{18, 21\}$

Step 2: Concatenate buckets (into a):

$\{0, 3, 6, 9, 12, 15, 18, 21\}$

Bucket Sort



Step 0: $a[0], \dots, a[n-1]$ are put in their corresponding bucket:

$w = b/m =$ bucket width

for $i=0, \dots, n-1$

Find bucket index: $j = a[i]/w$

Put $a[i]$ in bucket j

Step 1: Sorting:

for $j=0, \dots, m-1$

Sort bucket j

Step 2: Concatenate buckets:

for $j=0, \dots, m-1$

Move content of bucket j to a

Bucket Sort

Expected running time:

Step 0: Put in buckets:
 $O(n)$

Step 1: Sort buckets (using MergeSort or HeapSort):
 $O(m \cdot n/m \log(n/m))$ (assuming $a[0..n-1]$ are distributed evenly)

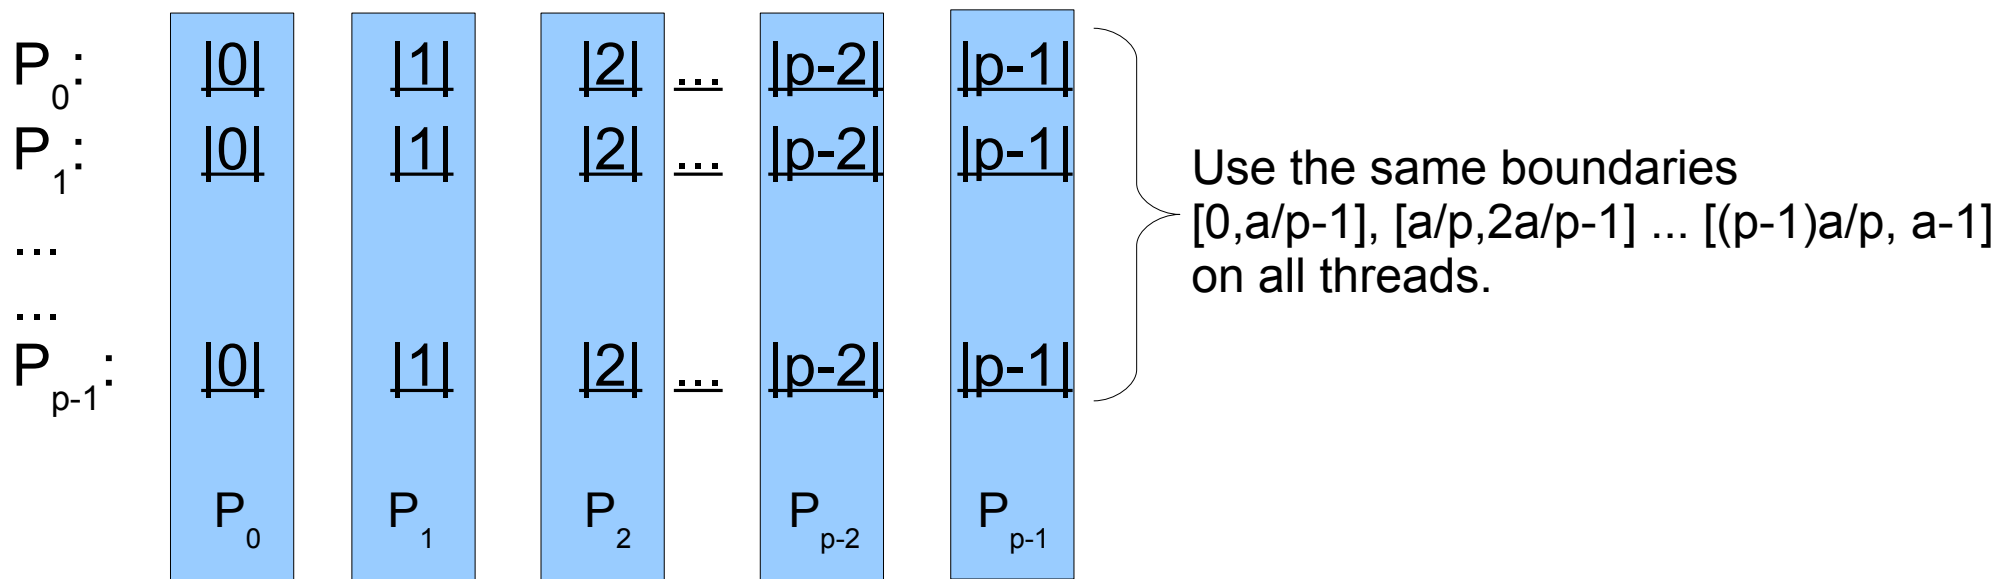
Step 2: Concatenate:
 $O(n)$ (assuming buckets are arrays)
 $O(m)$ if buckets are lists

Total time:
 $O(n \log(n/m))$
linear if $n/m=k$ is a constant: Choose $m=n/k$.

Parallel Bucket Sort

Algorithm:

- Divide data evenly between the threads,
- Put data into one of p local buckets.



- Thread P_i sorts the data in all the i 'th buckets

Issues:

Which local sorting algorithm to use? How much space is needed for buckets?
Where should each thread store its final result?

Generating random numbers

Pseudorandom numbers in parallel programs:

- Centralized: A dedicated thread generates all pseudorandom numbers
 - May involve many requests to the generating thread
 - Not very suitable for Monte Carlo simulation
- Distributed: Each thread generates the numbers as they are needed
 - Sequences must be distinct
 - Seeds must be distinct
 - Can use a hash function of (time and) process rank

`srandom()`, `random()`: Sequential methods for seeding and generating random numbers, all numbers are generated from the same sequence.

`srandom_r()`, `random_r()`: Thread safe, generates different sequences on different threads. Slightly more complicated to use.